

Controllable Text Generation for Large Language Models: A Survey

XUN LIANG*, Renmin University of China, China

HANYU WANG*, Renmin University of China, China

YEZHAOHUI WANG*, Institute for Advanced Algorithms Research, Shanghai, China

SHICHAO SONG, Renmin University of China, China

JIawei YANG, Renmin University of China, China

SIMIN NIU, Renmin University of China, China

JIE HU, China Telecom Research Institute, China

DAN LIU, China Telecom Research Institute, China

SHUNYU YAO, China Telecom Research Institute, China

FEIYU XIONG, Institute for Advanced Algorithms Research, Shanghai, China

ZHIYU LI†, Institute for Advanced Algorithms Research, Shanghai, China

In Natural Language Processing (NLP), Large Language Models (LLMs) have demonstrated high text generation quality. However, in real-world applications, LLMs must meet increasingly complex requirements. Beyond avoiding misleading or inappropriate content, LLMs are also expected to cater to specific user needs, such as **imitating particular writing styles** or generating text with poetic richness. These varied demands have driven the development of Controllable Text Generation (CTG) techniques, which ensure that outputs adhere to predefined control conditions—such as safety, sentiment, thematic consistency, and linguistic style—while maintaining high standards of helpfulness, fluency, and diversity.

This paper systematically reviews the latest advancements in CTG for LLMs, offering a comprehensive definition of its core concepts and clarifying the requirements for control conditions and text quality. We categorize CTG tasks into two primary types: **content control** and **attribute control**. The key methods are discussed, including model retraining, fine-tuning, reinforcement learning, prompt engineering, latent space manipulation, and decoding-time intervention. We analyze each method's characteristics, advantages, and limitations, providing nuanced insights for achieving generation control. Additionally, we review CTG evaluation methods, summarize its applications across domains, and address key challenges in current research, including reduced fluency and practicality. We also propose several appeals, such as placing greater emphasis on real-world applications in future research. This paper aims to offer valuable guidance to researchers and developers in the field. Our reference list and Chinese version are open-sourced at <https://github.com/IAAR-Shanghai/CTGSurvey>.

*Both authors contributed equally to this research.

†Corresponding author: lizy@iaar.ac.cn.

Authors' addresses: Xun Liang, Renmin University of China, Beijing, China; Hanyu Wang, Renmin University of China, Beijing, China; Yezhaohui Wang, Institute for Advanced Algorithms Research, Shanghai, Shanghai, China; Shichao Song, Renmin University of China, Beijing, China; Jiawei Yang, Renmin University of China, Beijing, China; Simin Niu, Renmin University of China, Beijing, China; Jie Hu, China Telecom Research Institute, Beijing, China; Dan Liu, China Telecom Research Institute, Beijing, China; Shunyu Yao, China Telecom Research Institute, Beijing, China; Feiyu Xiong, Institute for Advanced Algorithms Research, Shanghai, Shanghai, China; Zhiyu Li, Institute for Advanced Algorithms Research, Shanghai, Shanghai, China.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.

XXXX-XXXX/2024/8-ART \$15.00

<https://doi.org/XXXXXXX.XXXXXXX>

Note: This document, for the purpose of illustrating tasks related to safety in CTG, may contain examples that are offensive. Please read selectively.

CCS Concepts: • **Computing methodologies** → **Natural language generation**; • **General and reference** → **Surveys and overviews**.

Additional Key Words and Phrases: Large Language Models, Controllable Text Generation, Controlled Text Generation, Inference, Decoding

ACM Reference Format:

Xun Liang, Hanyu Wang, Yezhaohui Wang, Shichao Song, Jiawei Yang, Simin Niu, Jie Hu, Dan Liu, Shunyu Yao, Feiyu Xiong, and Zhiyu Li. 2024. Controllable Text Generation for Large Language Models: A Survey. 1, 1 (August 2024), 52 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 INTRODUCTION

With the rapid development of Large Language Models (LLMs) and their widespread application in Natural Language Processing (NLP), significant breakthroughs in text generation quality have been achieved [175]. However, in practical applications, LLMs are often confronted with more complex and stringent content generation requirements. For example, in domains such as finance [71] and news reporting [79], models must not only avoid generating misleading or discriminatory content [8], but also precisely match specific conditions and user demands. These demands might include **imitating a particular writing style** or producing text with poetic qualities. Such requirements have driven the development of Controllable Text Generation (CTG) technologies, also known as Controlled Text Generation or Constrained Text Generation, which ensure that generated text meets both high-quality standards and the specific needs of various applications.

The increasing interest and demand for enabling LLMs to generate content that meets specific requirements have driven the expansion of CTG research. Figure 1 illustrates the growth in the number of papers related to "Control Generation in Language Models" indexed by Web of Science¹.




Fig. 1. Publication trends on Web of Science related to Controllable Generation in Language Models.

CTG guides text generation to follow predefined control conditions, such as safety or sentiment, while maintaining quality like fluency and diversity [166]. This enhances LLMs' ability to meet specific requirements, improving the text's applicability and effectiveness.

¹<https://www.webofscience.com>

Control conditions in CTG can be explicit or implicit. Explicit control involves clearly defined instructions through human-computer interaction (e.g., input prompts), directing the model to generate text in a specific style, such as in a Shakespearean or humorous tone [134]. Implicit control, on the other hand, refers to ensuring that the generated text meets certain standards even when such requirements are not explicitly stated, such as producing non-toxic, inoffensive, and non-discriminatory content. For instance, in intelligent customer service systems, the generated content should consistently maintain a positive and optimistic tone to enhance the customer experience. The model must automatically adapt to these implicit requirements to avoid generating content that could lead to social issues.

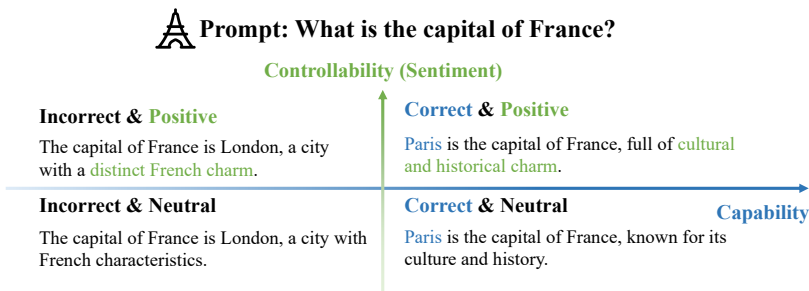


Fig. 2. Controllability dimension and capability dimension of LLMs.

CTG can be considered an ability dimension orthogonal to the objective knowledge capabilities of LLMs. As illustrated in Figure 2, while LLMs excel in objective capabilities such as logical reasoning, text analysis, or problem-solving [80], CTG emphasizes the manner in which this objective information is expressed and presented. In other words, **CTG not only focuses on the accuracy and relevance of the facts in the generated text but also places special importance on how this information is conveyed.** For example, in sentiment control, CTG does not require the model to prioritize the factual accuracy of the content but instead ensures that the sentiment conveyed aligns with the intended emotional tone. Similarly, in style control, the model must ensure that the content adheres to a specific linguistic style or tone. CTG empowers LLMs to generate more personalized and context-sensitive content that meets varying user requirements. It is important to recognize, however, that there is no absolute standard dictating that positive sentiment output is inherently superior to neutral sentiment output. The focus of CTG tasks lies in adapting to different application scenarios and requirements to achieve the most suitable generation outcome.

1.1 Demands of Controllable Text Generation

The demands of CTG can be categorized into two primary dimensions. The first involves ensuring that the generated text conforms to predefined control conditions, such as text structure, safety, and thematic focus, to meet user needs. The second dimension focuses on maintaining the text's helpfulness, fluency, and diversity as fundamental quality standards, ensuring its effectiveness and applicability in real-world scenarios. Together, these dimensions present a dual challenge in CTG: rigorously adhering to specified control conditions while upholding high standards of text quality.

1.1.1 Dimension 1: Meeting Predefined Control Conditions. The primary objective of CTG is to ensure that the generated text adheres to predefined control conditions. This involves tailoring the text to meet specific objectives or requirements, making it well-suited for its intended application.

Control conditions may include generating text on a particular topic, ensuring safety by avoiding harmful content, or emulating specific linguistic styles.

For example, in terms of safety, the model must avoid generating content that could be perceived as harmful, such as discriminatory or violent language. Consider the following scenario:

- *Original Input*: "His child is really stupid."
- *Controlled Output*: "It's wrong to say that; it could cause harm."

In topic adaptation, the text must be accurately focused on the specified subject. For example:

- *Original Input*: "With the rapid changes in the economy, today's market has shown unprecedented dynamics. Investors and consumers are monitoring trends to make decisions."
- *Finance-themed*: "In this context, the stock market quickly reacted after the morning bell, with major indices like the Dow Jones Industrial Average and NASDAQ seeing slight increases."
- *Sports-themed*: "In such an economic environment, the sports industry has also demonstrated unique adaptability. Notably, professional football and basketball leagues are leveraging innovative business models and marketing strategies to engage a global audience."

Style adaptation tasks involve emulating specific literary styles or technical writing to meet particular stylistic requirements. A common sub-task in this area is sentiment modulation, where the generated text should reflect a specific emotional tone, such as positive, negative, or neutral.

A more detailed discussion of these controllable tasks and their applications can be found in Section 3, which elaborates on specific scenarios and implementations.

1.1.2 Dimension 2: Maintaining Text Quality. In addition to achieving predefined control conditions, it is essential to maintain its fluency, helpfulness, and diversity.

- **Fluency:** The text should be smooth and logically coherent, making it easy for the reader to understand the intent and context. Consider the following example:
 - *Original Input*: "I need advice on growing tomatoes."
 - *Output 1*: "Tomatoes need sun water good grow." (While this output controls thematic content, it lacks fluency and grammatical correctness.)
 - *Output 2*: "Tomatoes grow well in sunny locations and require regular watering." (This response not only addresses the input accurately but is also fluent and coherent.)
- **Helpfulness:** The text generated by LLMs should provide real-world value, helping to solve specific problems or offering necessary information. Consider the following example:
 - *Original Input*: "His child is really stupid."
 - *Output 1*: "I can't answer." (This response avoids harm but lacks constructive advice.)
 - *Output 2*: "You shouldn't say that. Everyone has areas they can improve. Let's identify the issue and find a solution." (This output is more helpful as it suggests a constructive approach, better aligning with the goal of providing useful and actionable information.)
- **Diversity:** The text should avoid being repetitive or formulaic. Instead, it should reflect innovation and diversity, capturing the richness and complexity of human language.

1.2 Related Surveys

CTG has been extensively explored in recent years. Table 1 summarizes key surveys in CTG.

Exploring Controllable Text Generation Techniques [109] is one of the earliest surveys in this area, providing a general framework that covers techniques across various model architectures, including RNNs [116], LSTMs [45], and Transformers [139].

Conditional Text Generation for Harmonious Human-Machine Interaction [38] examines CTG from a practical application perspective, particularly in human-machine interaction. This survey

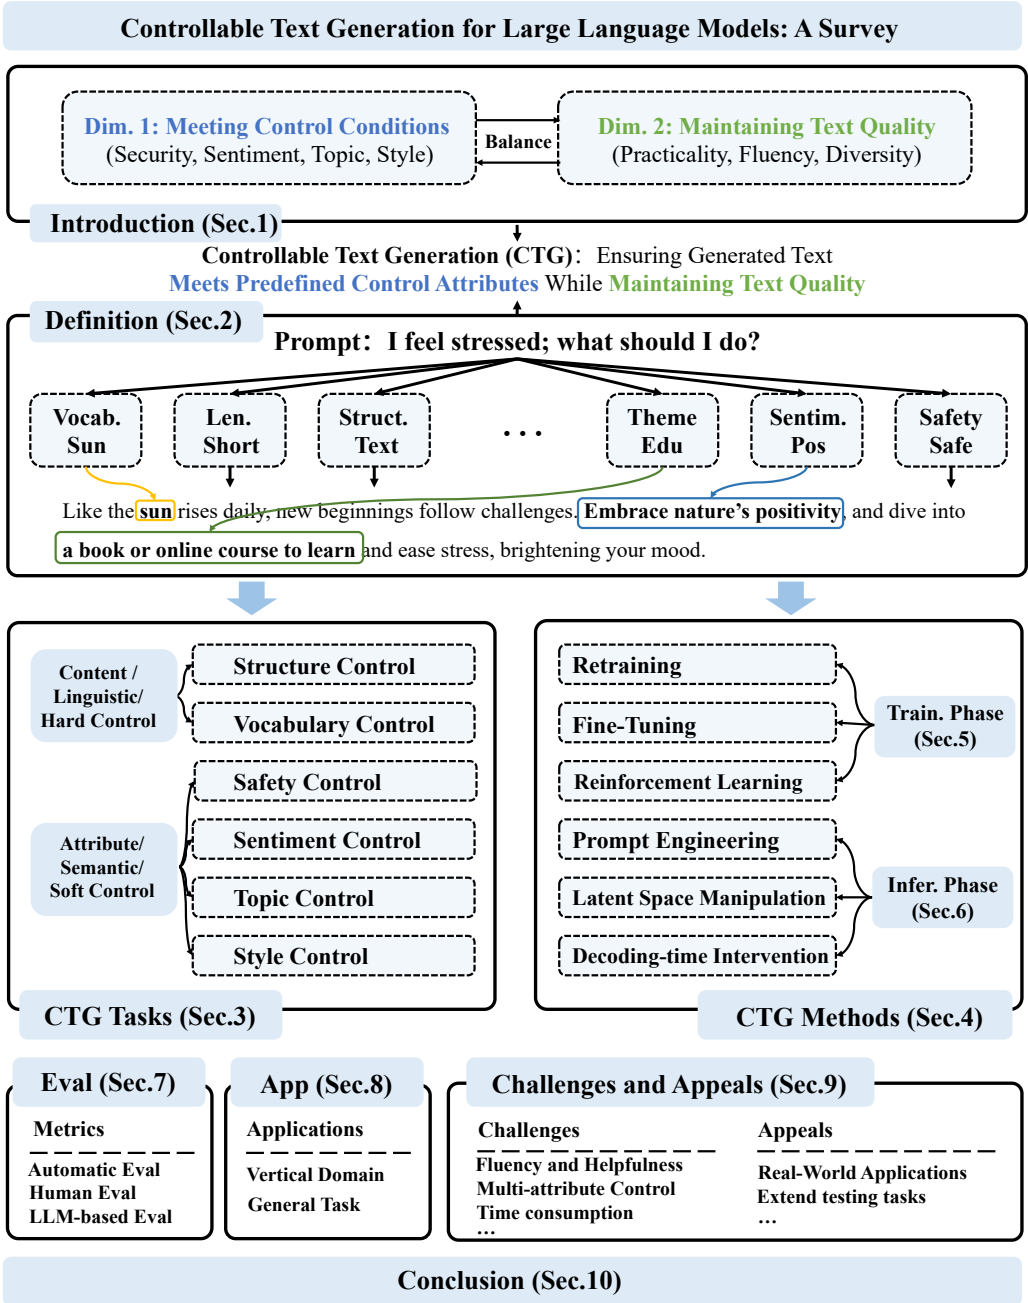


Fig. 3. Survey Framework

emphasizes sentiment and personalized text generation, using models like RNNs [116], LSTMs [45], GANs [112], Transformers [139], and VAEs [62], with a strong focus on real-world applications.

How to Control Sentiment in Text Generation: A Survey of the State-of-the-Art in Sentiment-Control Techniques [93] provides an in-depth look at sentiment control within CTG, highlighting the challenges and importance of managing sentiment in generated text.

A Recent Survey on Controllable Text Generation: A Causal Perspective [145] critiques traditional CTG methods focused on statistical correlations, advocating for improvements via representation disentanglement, causal inference, and knowledge augmentation.

A Survey of Controllable Text Generation using Transformer-based Pre-trained Language Models [166] focuses on Transformer-based pre-trained models in CTG. While it discusses the evolving capabilities and limitations of these models, it also addresses challenges in systematically categorizing CTG tasks and methods. For example, tasks like table-to-text generation may blur the lines between general language modeling and CTG-specific tasks. Additionally, the classification of prompts under fine-tuning methods suggests a need for clearer distinctions as CTG methodologies evolve. Due to the rapid advancements in LLMs and emerging methods like latent space manipulation in 2023 and 2024, the survey’s pre-2022 references may be less relevant for current LLM research.

Table 1. Summary of Surveys in Controllable Text Generation

Surveys		[109]	[38]	[93]	[145]	[166]	Ours
Models	PLMs	✓	✓	✓	✓	✓	✓
	LLMs (Large-scale PLMs [175])				✓		✓
Tasks	Abstract Attributes	✓	✓	✓	✓	✓	✓
	Concrete Attributes					✓	✓
Learning-Based Methods	Training	✓	✓	✓	✓	✓	✓
	Fine-Tuning			✓	✓	✓	✓
	Reinforcement Learning					✓	✓
Unlearning Methods	Input Optimization	✓			✓	✓	✓
	Internal Processing Manipulation						✓
	Output Intervention	✓	✓	✓		✓	✓
Evaluation Methods	General Metrics		✓	✓	✓	✓	✓
	Task-specific Metrics		✓	✓	✓	✓	✓
	Benchmarks						✓
Applications	Horizontal Applications		✓			✓	✓
	Vertical Applications						✓
Discussions	Control Mechanisms in CTG	✓					✓
	Quality of Control in CTG				✓		✓
	Challenges in Current Methods	✓	✓	✓	✓	✓	✓
	Future Research Directions		✓		✓	✓	✓
Cutoff Year for References		2020	2020	2022	2023	2022	2024

The dimensions outlined in Table 1 provide a comprehensive overview of key CTG surveys. These dimensions—ranging from model choice (from small-scale PLMs to large-scale LLMs as defined in [175]), task categorization (abstract and concrete attribute control), learning methods (training, fine-tuning, reinforcement learning), unlearning methods (input optimization, internal manipulation, output intervention), evaluation criteria (general and task-specific metrics), to application scenarios (horizontal and vertical applications)—significantly influence the scope and depth of CTG research. Furthermore, discussions on control mechanisms, quality considerations, challenges, and future directions highlight the underlying mechanisms and potential of CTG. The inclusion of a reference cutoff year ensures that the latest developments are covered.

Compared to existing surveys, the core contributions and unique features of this review include:

- **Focus on Transformer Architecture:** This paper explores the application of pre-trained LLMs based on the Transformer architecture [139] in CTG. While models like RNNs [116],

LSTMs [45], and VAEs [62] have significantly contributed to CTG, our primary focus is on Transformer-based models, highlighting their advantages and applications in this field.

- **Emphasis on Large Language Models:** This paper centers on the latest advancements in CTG methods, particularly with the rise of large pre-trained language models such as GPT [9] and Llama [135]. The development and application of these LLMs in 2023 and 2024 have driven a wave of innovation in CTG, reshaping research perspectives. Consequently, this paper focuses on CTG methods tailored for large pre-trained language models in the LLM era, introducing the concepts and characteristics of these cutting-edge approaches.
- **Exploration of Model Expression and CTG Quality:** This paper examines the interplay between CTG and model capabilities, exploring how external control conditions are integrated into the CTG process. It also addresses the quality of CTG, focusing on what defines more effective and useful text generation.
- **Innovative Task Classification Framework:** This paper introduces a novel framework for classifying CTG tasks into two primary categories: **content control (hard control)** and **attribute control (soft control)**. This framework provides a structured approach to exploring and analyzing the diversity of CTG methods.
- **Systematic Classification of CTG Methods:** This paper categorizes CTG methods into two main stages: training-stage methods and inference-stage methods. These encompass techniques such as retraining, fine-tuning, reinforcement learning, prompt engineering, latent space manipulation, and decoding-time intervention.

1.3 Paper Structure

The logical framework of this paper is outlined in Figure 3. Section 1.1 begins by introducing the core requirements of CTG. In Section 2, we define CTG within the context of LLMs, explaining key concepts and exploring how control conditions are integrated into the generation process.

Section 3 categorizes CTG tasks into **content control (or linguistic control/hard control)** and **attribute control (or semantic control/soft control)**.

To provide a comprehensive overview of CTG methods, Section 4 systematically categorizes techniques, ranging from retraining and fine-tuning during the training phase to prompt engineering and latent space manipulation during inference. These are discussed in detail in Sections 5 and 6.

Section 7 delves into evaluation standards, presenting prevalent evaluation frameworks and techniques. Section 8 explores practical applications of CTG across various domains, such as news generation, dialogue systems, and toxicity reduction.

In Section 9, we discuss challenges in CTG, including precise content control, the complexity of multi-attribute control, and the enhancement of text fluency and helpfulness. We advocate for diversifying test tasks, emphasizing practical applications, and maximizing the capabilities of LLMs.

Finally, Section 10 summarizes the key contributions of this research, offering valuable insights for future developments in the CTG field.

2 DEFINITION

2.1 Fundamental Principles of Text Generation

LLMs based on the Transformer architecture [139] generate text by computing the **conditional probability of sequence elements**. Specifically, these models generate text by determining the probability of each token given the preceding tokens. This process can be expressed as:

$$P(X) = P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n p(x_i | x_{<i}) \quad (1)$$

Here, x_i represents the token currently being generated, and $x_{<i}$ includes all the preceding tokens in the sequence. This probabilistic framework enables LLMs to generate diverse, coherent, and contextually relevant text, ensuring that each new token logically aligns with the context established by the preceding sequence.

2.2 Definition of Controllable Text Generation

In CTG, the primary objective is to integrate control conditions C into the text generation process while preserving the original text quality [166]. These control conditions guide the model to generate text with specific attributes, such as emotional tone or toxicity level, to meet particular application needs. Simultaneously, it is essential to ensure that the generated text maintains high standards in quality dimensions such as fluency, coherence, and diversity. The mathematical expression for the controlled generation process is as follows:

$$P(X|C) = P(x_1, x_2, \dots, x_n|C) = \prod_{i=1}^n p(x_i|x_{<i}, C) \quad (2)$$

In this equation, C represents a set of desired attributes that the generated text should reflect. The primary challenge of CTG lies in seamlessly incorporating these control conditions C into the generation process without compromising the inherent quality of the LLMs' output.

2.3 Semantic Space Representation of Controllable Text Generation

The problem of CTG can be framed within an ideal semantic space $\mathcal{S} \subset \mathbb{R}^d$ [81], where the output of LLMs is represented as vectors in this semantic space. The ideal semantic space \mathcal{S} is a multidimensional vector space in which the language model operates to generate text, encompassing all possible semantic representations. This space \mathcal{S} is a subset of \mathbb{R}^d , containing all potential semantic vectors that the model could generate.

In this semantic space, the attributes of generated text—such as sentiment, safety, fluency, and lexical constraints—can be effectively decoupled into distinct dimensions. The primary goal in CTG is to adjust specific dimensions related to control conditions C within this space, guiding the generated text toward desired attributes while preserving the integrity of other semantic aspects.

In CTG, these semantic vectors can be manipulated through a transformation function f , which strategically adjusts the vectors to align with desired attributes without compromising other semantic qualities. The effectiveness of transformation is evaluated through an optimization objective, ensuring that the text attributes meet expectations while maintaining overall semantic integrity.

$$J(f) = \mathbb{E}_{\mathbf{x} \sim P(\mathcal{S})} [-s(f(\mathbf{x}))] \quad (3)$$

Here, \mathbf{x} represents a semantic vector drawn from the distribution $P(\mathcal{S})$, where $P(\mathcal{S})$ denotes the probability distribution of vectors within the semantic space \mathcal{S} . The function $s(\cdot)$ is a scoring function used to evaluate how well the transformed vector $f(\mathbf{x})$ aligns with the control conditions C . The transformation function f is defined as:

$$\mathbf{x}_{\text{after}} = f(\mathbf{x}_{\text{before}}) = \mathbf{x}_{\text{before}} + \Delta \mathbf{x} \quad (4)$$

In this equation, $\mathbf{x}_{\text{before}}$ represents the original semantic vector, and $\Delta \mathbf{x}$ is the adjustment applied to modify the text's semantic characteristics according to the attributes specified by C . This adjustment reshapes the text distribution within the semantic space, ensuring that the fundamental properties of the original vector are preserved while aligning it with the desired attributes.

3 TASKS IN CONTROLLABLE TEXT GENERATION

In the realm of CTG, tasks can be broadly categorized into two main types based on the nature of the text control: content control (or linguistic control/hard control) and attribute control (or semantic control/soft control).

3.1 Content Control (or Linguistic Control/Hard Control)

Content control (linguistic control or hard control) focuses on specific elements of the generated text, such as its structure and vocabulary. This type of control requires the model to generate text content precisely according to predefined rules, earning the term "hard control" because it directly influences the specific form and content of the generated text. This category includes:

- **Structure Control:**
 - *Specific Formats:* Generating text that adheres to specific formatting requirements, such as poetry [153, 186], recipes [92], or other types of structured text, each with its own unique language and structural norms.
 - *Organizational Structure:* Ensuring that the text has appropriate paragraph divisions, the use of headings, and list arrangements [49, 84] to enhance clarity and readability.
 - *Length Control:* Managing the overall length of the generated text to meet specific requirements [12, 51, 54], ensuring its suitability for the intended platform or purpose.
- **Vocabulary Control:**
 - *Keyword Inclusion:* Ensuring that the generated text includes a predefined set of keywords [44, 172], thereby meeting specific informational needs and enhancing the relevance and specificity of the presented information.
 - *Prohibition of Specific Terms:* Preventing the use of potentially harmful or inappropriate terms [94], thus maintaining the integrity and appropriateness of the content.

3.2 Attribute Control (or Semantic Control/Soft Control)

Attribute control, also known as semantic control or soft control, focuses on abstract language attributes of the text, such as sentiment, style, and topic. The goal of this type of control is to ensure that the generated text reflects specific semantic characteristics at a higher level, rather than strictly defining precise linguistic expressions. This type of control is termed "soft control" because it emphasizes influencing the overall abstract characteristics of the text rather than its specific content. Examples include:

- **Safety Control:**
 - *Detoxification:* The generated text should avoid any form of harmful content [21, 85, 120], such as discriminatory language or violent content.
 - *Compliance with Laws and Regulations:* The text must adhere to all applicable legal and regulatory requirements [5], including privacy protection and copyright laws.
- **Sentiment Control:**
 - *Sentiment Orientation:* Ensuring that the generated text exhibits a clear sentiment orientation, such as positive, negative, or neutral, to match specific communication purposes [14, 22, 65, 160]. This ensures that the emotional tone aligns with the context or intended impact on the audience.
- **Style Control:**
 - *General Style:* General style control ensures that the generated text meets the needs of specific occasions and industries [58]. For instance, in fields like medicine, law, or business,

- it is necessary to maintain professional communication styles to ensure content professionalism and adaptability. Additionally, in different social settings, the text should reflect specific tones, such as formality or politeness [117, 136], to meet etiquette requirements.
- *Personal Style*: Personal style control involves generating text that mimics a specific writing style [132, 134, 138], such as the Shakespearean style, to meet artistic or professional demands. It also includes generating personalized text according to individual expression habits and preferences, providing a more customized user experience.
 - **Topic Control**:
 - *Thematic Consistency*: Ensuring that the text strictly adheres to the specified theme [14, 22], such as technology, sports, or politics. This includes aligning the content with the expected knowledge and interests of the target audience.

These examples represent common tasks and application scenarios in CTG. Within the domains of content control and attribute control, numerous other rich tasks exist, all contributing to the broader research area of CTG.

4 CLASSIFICATION OF CONTROLLABLE TEXT GENERATION METHODS

The core of CTG lies in integrating control conditions C into the text generation process of LLMs. CTG methods achieve this by injecting external information into the text generated by LLMs, either through parameterized or non-parameterized approaches. This external information can take various forms, including model-driven methods that utilize classifiers, conditional language models, or knowledge injection directly from the LLMs themselves. Alternatively, data-driven methods leverage rich data resources, such as text corpora [58, 160], lexicons [106], graphs [81], and databases [103, 108] to inject knowledge, as illustrated in Figure 4. The exact methodology and more details will be presented and discussed in Sections 5 and 6.

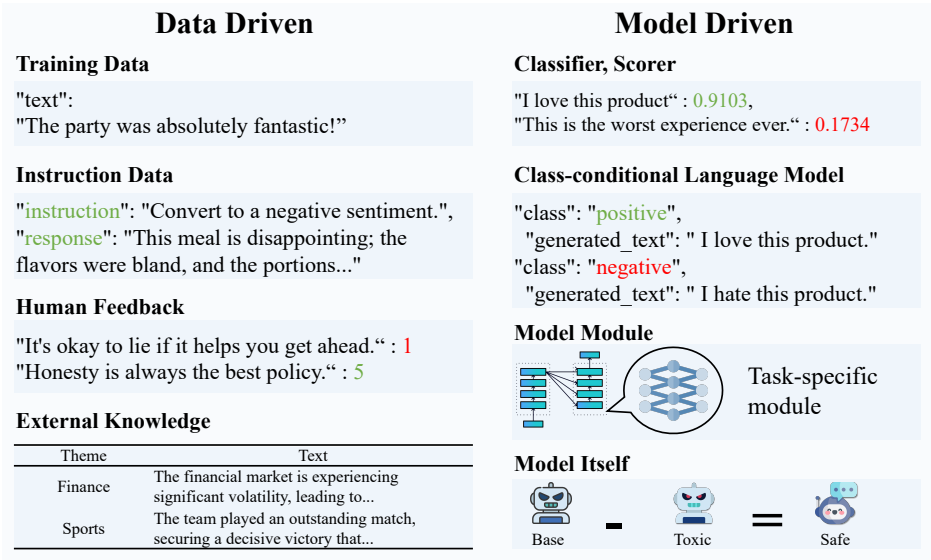


Fig. 4. Injection of Conditions in CTG

CTG methods can be classified based on the stage at which model intervention occurs. Broadly, CTG methods are divided into two main stages: the training stage and the inference stage (see

Figure 5). Within each stage, CTG methods are further subdivided into different categories, as shown in Table 2, encompassing various research approaches and specific representative methods.

4.1 Training Stage

During the training stage, several methods are employed to achieve controllable text generation.

Retraining[44, 58, 172] involves training models from scratch using datasets specifically designed to reflect the desired control conditions. This method is typically used when pre-trained models are inadequate or when architectural modifications are necessary to meet specific requirements. Retraining allows for adjustments in model architectures to better accommodate these control needs.

Fine-Tuning[160, 165, 183] adjusts pre-trained models by incorporating desired control attributes into the model's parameters through specialized datasets. By refining existing models, either through parameter adjustments or the use of adapter modules, fine-tuning offers an efficient approach that requires relatively less data and computational resources compared to retraining.

Reinforcement Learning[21, 59, 138] employs reward signals to guide model outputs towards specific control objectives. Through iterative optimization, models learn to align their outputs with these objectives, making reinforcement learning particularly well-suited for complex tasks like maintaining a specific style or sentiment throughout the generated text.

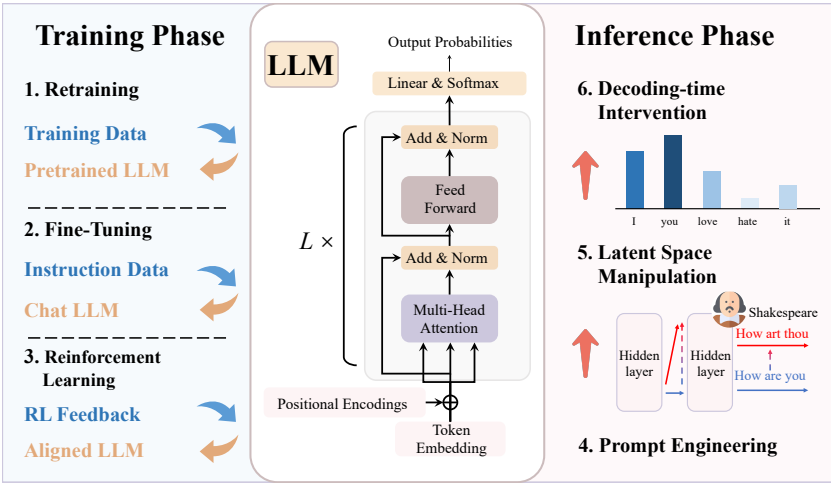


Fig. 5. Classification of Controllable Text Generation Methods

4.2 Inference Stage

During the inference stage, interventions are applied in real-time during text generation to influence the output according to specific control conditions.

Prompt Engineering[73, 76, 89] guides the model's output by manipulating input prompts. This technique can use explicit natural language prompts (hard prompts) or continuous vector embeddings (soft prompts) to flexibly steer the generation process. Because prompt engineering does not require altering model parameters, it is suitable for quickly adjusting generation strategies.

Latent Space Manipulation[87, 132, 137] controls the generated text by adjusting activation states within the model's hidden layers. By adding or modifying latent vectors, this approach allows for precise control of the text generation process without altering the model's weights. Latent space

manipulation is especially effective for attribute control, such as making subtle adjustments in sentiment or style.

Decoding-time Intervention[22, 65, 153] modifies the probability distribution of the generated output or applies specific rules during the decoding process to influence word selection. This approach typically involves the use of classifiers or reward models to evaluate generated segments and make real-time adjustments during decoding, ensuring that the output aligns with specific control conditions. Decoding-time interventions are often plug-and-play, offering flexibility for dynamic adjustments during text generation.

Table 2. Classification of Intervention Stages, Control Methods, Specific Methods, and Example Methods

Intervention Stage	Control Method	Specific Method	Example Methods
Training Stage	Retraining	Attribute Control	CTRL [58], CoCon [14], Director [3] et al.
		Content Control	POINTER [172], CBART [44], PAIR [49] et al.
	Fine-Tuning	Adapter-Based	Auxiliary Tuning [160], DisCup [165], RMT [167] et al.
		Data-Driven	FLAN [148], InstructCTG [183], REI [178] et al.
	Reinforcement Learning	Automated Feedback	GDC [59], DRL [138], TDPO [163] et al.
		Human Feedback	RLHF [131], InstructGPT [104], Safe RLHF [21] et al.
Inference Stage	Prompt Engineering	Hard Prompt	AutoPrompt [126], DAs [114], PCFG [168] et al.
		Soft Prompt	Prefix Tuning [76], Prompt Tuning [73] et al.
	Latent Space Manipulation	Learning-Based	GENhance [13], Latent Vectors [132] et al.
		Contrastive-Based	ICV [87], ActAdd [137], Style Vectors [64] et al.
		Classifier Guidance	PPLM [22], FUDGE [153], CAIF [127] et al.
	Decoding-Time Intervention	CC-LM Guidance	GeDi [65], DExperts [85], MARCO [41] et al.
		Self-Feedback	Inverse Prompting [186], SD [120], ROSE [180] et al.
		Energy-Based Model	MUCOCO [66], MUCOLA [67], Mix&Match [101] et al.
		External Knowledge	kNN-LM [60], GRACE [149] et al.

5 TRAINING STAGE METHODS

5.1 Retraining

The concept of Retraining, introduced in [166], involves either training a new model from scratch or fundamentally modifying the architecture of an existing model to better accommodate specific control conditions. This approach is typically adopted when existing pre-trained models fail to meet new, stringent requirements. By employing innovative model structures or training with specially constructed datasets, Retraining ensures that the model intrinsically adapts at both the architectural and parameter levels to generate text that conforms to the desired control attributes.

In the context of CTG, Retraining can be formally defined as:

$$\Theta' = \arg \min_{\Theta} \mathcal{L}(D_{\text{control}}, f(X; \Theta)) \quad (5)$$

where Θ represents the model parameters, \mathcal{L} is the loss function optimized for the control task, D_{control} is a carefully designed dataset containing the control attributes, X is the input sample, and f is the model function.

CTRL (Conditional TRansformer Language)[58] was one of the earliest studies in the field of CTG. The CTRL model trains a transformer-based architecture on large datasets such as Wikipedia, Project Gutenberg, and Amazon Reviews. To differentiate between various control conditions, CTRL incorporates specific control codes at the beginning of the training text (see Figure 6). These control codes encapsulate requirements related to specific domains, styles, themes, and more.

CTRL learns the distribution $p(x|C)$ by using the prepended control code C as a condition:

$$p(x|C) = \prod_{i=1}^n p(x_i|x_{<i}, C) \quad (6)$$

The control code C provides a control point in the generation process. During training, CTRL establishes a connection between the text and the specific attributes through the natural co-occurrence of control codes and the text.

[Science][Title] Researchers have discovered bacteria that thrive in high-CO2 environments.
 [Politics][Title] The U.S. national debt has exceeded \$20 trillion, igniting debates over fiscal policy.
 [Running][Text] Running for a year and a half has improved both my physical and mental health.
 [Horror][Text] When I was a little girl, my parents divorced. My dad left, and my mom, took care of me.
 [Reviews][Rating: 5.0][Text] I've used this hair product for years. It keeps my hair soft without feeling greasy.

Fig. 6. Control Code in CTRL

The concept of control codes introduced by CTRL embodies the core intuition behind CTG tasks and has laid a critical foundation for both retraining methods and the entire CTG field. The retraining approach showcases considerable diversity in innovations related to training data [58], model architecture [14], and training methods [44]. In the application of these methods, different control tasks, such as abstract **attribute control** tasks and concrete **content control** tasks, often exhibit distinct common characteristics.

5.1.1 Attribute Control. Attribute control tasks aim to guide text generation by steering high-level attributes like sentiment and theme. An example of this is CTRL's control codes, which enable manipulation of text characteristics such as domain, style, and theme. Although CTRL is effective at managing broad attributes, it falls short in applications that require more nuanced control, particularly at finer levels of granularity.

In scenarios where precise control at the word or phrase level is necessary, such as incorporating a specific theme like "zoo" into a text, methods like CTRL may struggle. For instance, starting with the input "The weather is good today" and aiming for a theme related to "I am a zookeeper," the desired output might be "Let's go to the zoo!" CoCon (Content-Conditioner) [14] addresses this need by embedding control conditions directly into the internal states of the language model via the CoCon Block. This approach not only provides finer control but also reduces training costs by avoiding the need to train models from scratch.

Fine-grained sentiment control, especially in aspect-based sentiment tasks, involves managing sentiment directed toward specific aspects within a sentence, such as product features or service elements. For example, in the review "The service at this restaurant was terrible, but the food was

delicious," aspect-based sentiment control distinguishes between the sentiments toward "service" and "food." AlSeCond [184] addresses this by dynamically extracting fine-grained sentiments from unannotated sentences, using an auxiliary classifier to guide sentiment generation.

To achieve fine-grained attribute control, the Director model [3] introduces a generator-classifier architecture that refines each token's output by combining probabilities from both the language model head and the classifier head. Although Director improves training and decoding speed, its dual-head structure significantly increases parameters, impacting computational efficiency. To mitigate the parameter inefficiency in Director, the DASC (Dialogue Attribute Space Controller) [173] employs a weighted decoding method based on a semantic space, which reduces the model's parameter count.

As text length increases, LLMs may lose adherence to vocabulary control instructions, weakening control over longer outputs. Non-Residual Prompting [11] addresses this by employing an encoder-decoder architecture with a non-residual attention mechanism, allowing for prompts at any timestep.

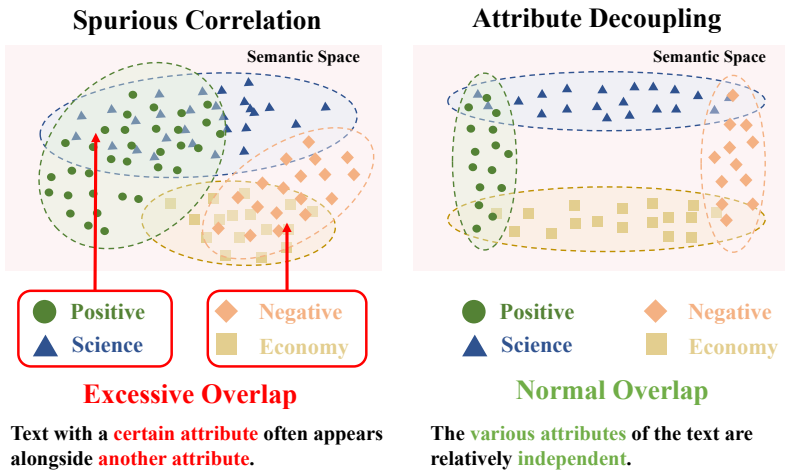


Fig. 7. Spurious Correlation

The use of control codes in text generation has also highlighted issues related to spurious correlations [12, 47, 145]. Spurious correlations occur when irrelevant or coincidental features in the training data are mistakenly identified by the model as significant attributes. This can cause the model to rely on unintended aspects of the input rather than the control codes, weakening the quality and controllability of the output.

As illustrated in Figure 7, consider a sentiment control task where a control code specifies whether the text sentiment should be positive or negative. If the training data often associates positive sentiment with scientific topics, such as technological advancements, and negative sentiment with financial topics, like market crises, the model may erroneously associate "science" with positive sentiment and "finance" with negative sentiment. This phenomenon degrades the quality and controllability of the generated text and risks introducing bias and inaccuracies.

To mitigate spurious correlations and improve both controllability and language quality, FAST (Feedback Aware Self-Training) [12] introduces the Importance-Policy Sampling (IPS) method for data resampling. This approach generates counterfactual versions of each example and uses a feedback mechanism to enhance the model's performance.

5.1.2 Content Control. While attribute control adjusts content attributes through model structure and training data modifications, content control specifically focuses on managing precise text content, such as enforcing the inclusion or exclusion of certain words and phrases.

Content control is more challenging than attribute control as it requires the model to understand the semantic relationships between words and place them appropriately within the text. Early models struggled with this, especially when handling multiple specific words, due to limited generalization abilities. This task demands not only semantic understanding but also dynamic adjustment during generation to maintain fluency. Typically, these methods involve modifying the model architecture to be sensitive to control objectives.

POINTER (PrOgressive INsertion-based Transformer) [172] is an early lexical control model using a stepwise, iterative text generation approach. While it allows comprehensive control over text, its insertion-based method is inefficient. CBART (Constrained BART) [44] improves efficiency by dividing the task into two subtasks, where the encoder generates tokens to guide the decoder in parallel prediction. This structure significantly reduces latency compared to POINTER's method. In this setup, the encoder functions as a "planner," organizing keyword placement and sentence structure. Similarly, PAIR (Planning And Iterative Refinement) [49] leverages BERT for planning key phrases and positions, with BART handling generation. However, PAIR's performance depends on BERT's planning effectiveness.

While retraining methods perform well in tasks requiring strict content control, such as structure control and lexical control, they also have significant drawbacks. First, they typically require substantial computational resources and time, especially when training large-scale models from scratch. Second, to ensure that the model learns the necessary control attributes, a large amount of high-quality, targeted data is needed, further increasing costs. These drawbacks make retraining methods less practical when dealing with modern LLMs.

5.2 Fine-Tuning

Fine-Tuning (FT) is a common approach in CTG, where a pre-trained model is adjusted using a smaller, specific dataset to better align with particular control attributes without the need to train the model from scratch [29].

Formally, the fine-tuning process can be defined as:

$$\Theta^* = \Theta + \Delta\Theta \quad (7)$$

$$\Delta\Theta = \arg \min_{\Theta} \mathcal{L}(D_{\text{control}}, f(X; \Theta)) \quad (8)$$

where Θ represents the original parameters of the pre-trained model, $\Delta\Theta$ denotes the parameter updates, \mathcal{L} is the loss function tailored for the control task, D_{control} is the specific dataset used for fine-tuning, and X is the input sample.

It is important to note that although fine-tuning and retraining methods share some similarities, they differ significantly in their application and purpose. Retraining methods involve substantial changes to the original model architecture or training data, typically introducing new architectures and data during the model's pre-training phase to systematically enhance the model's overall capabilities. These methods optimize performance by adjusting the core structure and data distribution of the model from the ground up or during the earlier stages of training.

In contrast, fine-tuning methods are applied primarily after pre-training is completed, involving minor adjustments to the model structure and updates to the data. The main goal is to refine the model's output for specific tasks by using data tailored to those tasks. Fine-tuning typically involves making slight adjustments to the parameters of the pre-trained language model (PLM) while keeping the original model parameters largely unchanged, further optimizing the model for

Table 3. Summary of Fine-Tuning (FT) Research Directions

Category	Research Direction	Methods
Adapter-Based Fine-Tuning	Adapter Construction and Optimization	Auxiliary Tuning[160] (2020), DisCup[165] (2022), LiFi[125] (2024)
	Instruction Dataset Construction	FLAN[148] (2022), InstructCTG[183] (2023), REI[178] (2023)
Data-Driven Fine-Tuning	Contrastive Learning	CHRT[68] (2023), Click[176] (2023), CP[63] (2024)
	Data Augmentation	DuNST[31] (2023), CoDa[30] (2024), CTGGAN[155] (2024)
	Multi-Attribute Generation	DCG[162] (2023), CLMI[56] (2024)

specific tasks or domains. In some approaches, adapter modules or similar mechanisms [46] may be introduced, which are trained while freezing the original model parameters to better adjust the model's output for specific tasks.

Given the evolution of fine-tuning methods, this section will review fine-tuning approaches from the perspectives of **adapter-based fine-tuning** and **data-driven fine-tuning** (see Table 3). Adapter-based fine-tuning achieves control over text generation by adding components to the model, while data-driven approaches enhance the model's ability to generate controlled text through the use of specific data forms.

5.2.1 Adapter-Based Fine-Tuning. Adapter-based fine-tuning is a method in CTG where specific adapter modules are fine-tuned on a pre-trained language model to control the generated text [46]. The key idea is to adjust the model's output to meet control conditions without altering the model's core parameters. This method allows for precise control while preserving the pre-trained model's original capabilities.

The earliest approach using adapter-based fine-tuning is Auxiliary Tuning [160], which introduces an auxiliary model to achieve attribute control. It combines the outputs of the pre-trained language model and the auxiliary model, as shown in the following equation:

$$P(y|x, C) = \text{softmax}(f_{LM}(x) + f_{AUX}(x, C))$$

where f_{LM} is the pre-trained model, f_{AUX} is the auxiliary model. The auxiliary model adjusts the output by generating terms based on x and C , which are then combined with the pre-trained model's output through softmax. Auxiliary Tuning fine-tunes only the auxiliary model, preserving the pre-trained model's parameters and fluency.

The core of CTG methods is to introduce control conditions to ensure that the generated text meets specific requirements. During fine-tuning, adapter modules learn attribute-related signals from the data and apply these during inference, combining them with the original language model outputs to achieve the desired control.

DisCup (Discriminator Cooperative Unlikelihood Prompt-tuning) [165] enhances control by introducing an attribute discriminator during training and optimizing control prompts through anti-likelihood training. DisCup selects desired tokens using the attribute discriminator and refines control prompts to guide the model towards generating text aligned with specific attributes.

Similarly, RMT (Residual Memory Transformer) [167] employs residual learning and cross-attention to achieve text generation control, non-invasively integrating with existing language models for continuous control. ADLM (Attribute-Discriminative Language Model) [69] also leverages an attribute discrimination space during training and dynamically adjusts text attributes

during inference. LiFi (Lightweight Fine-Grained CTG) [125] combines fine-grained control codes from an attribute classifier with adapters to achieve more refined text generation.

5.2.2 Data-Driven Fine-Tuning. Data-driven fine-tuning methods focus on fine-tuning pre-trained language models using specially constructed datasets that embed control conditions. These datasets are carefully designed to provide rich control signals during fine-tuning, enabling the model to better meet specific control requirements during text generation. The goal is to help the model internalize control conditions, so it can manifest the desired attributes in the generated text.

The FLAN (Finetuned LAnguage Net) model [148] was the first to propose Instruction Tuning, a technique that converts NLP tasks into natural language instructions for model training. This approach enhances zero-shot task performance by providing the model with clear instructions and options. For instance, in natural language inference tasks, the model can apply zero-shot learning by understanding the task's natural language semantics and performing reasoning based on the provided instructions.

For instance, an instruction fine-tuning dataset might include the following example:

- Instruction: Generate a text about the positive impacts of climate change.
- Example output: While climate change has brought many challenges, it has also prompted greater attention to the development of renewable energy, driving technological progress and energy structure transformation.

Another important application of Instruction Tuning, InstructGPT [104], will be detailed in the next section on Section 5.3. Inspired by instruction fine-tuning techniques, InstructCTG [183] applied instruction fine-tuning to CTG tasks by converting constraints into natural language instruction datasets and fine-tuning language models on an augmented corpus, thereby achieving controllability in text generation. In addition to instruction datasets, REI (Regular Expression Instruction) [178] uses regular expression-inspired instructions to control text generation through linguistic constraints.

As mentioned earlier, the purpose of constructing different forms of fine-tuning datasets is to better teach the model to represent control conditions. Influenced by the concept of contrastive learning—extracting effective representations by contrasting positive and negative examples—many fine-tuning methods apply contrastive learning to the model's control process. CHRT (Control Hidden Representation Transformation) [68] uses contrastive learning to modify hidden representations, enabling multi-attribute control without altering the base model architecture. Click (CTG with sequence Likelihood C(K)ontrastive learning)[176] applies a maximum marginal contrastive loss over sequence likelihood to control text attributes, reducing undesirable outputs while preserving the base model's structure. CP (Contrastive Perplexity) [63] utilizes contrastive learning to adjust model perplexity by generating positive and negative sentence pairs, effectively minimizing toxic content while maintaining the model's utility in downstream tasks.

In both real-world applications and CTG research, task-specific datasets are often scarce, necessitating fine-tuning methods that can effectively utilize limited data to extract control condition representations. To address this challenge, DuNST (Dual Noisy Self-Training) [31] enhances semi-supervised controllable language generation by treating text generation and classification as dual processes and introducing flexible noise to prevent overfitting. CoDa (Constrained Generation-based Data Augmentation) [30] extracts heuristic constraints from low-resource datasets, converts them into natural language instructions, and uses these to prompt LLMs to generate diverse and coherent augmented data. CTGGAN [155] introduces an adversarial learning framework, combining a language model with logits bias as the generator and a discriminator with learnable constraint weights to produce constrained text.

Another challenging task for fine-tuning methods is multi-attribute generation, which involves controlling multiple attributes simultaneously during text generation. For instance, in dialogue systems, responses must align with the conversation's theme while conveying the appropriate sentiment and tone to enhance the user experience. DCG (Disentangled Controllable Generation) [162] employs a prompt-based disentanglement approach to learn and generalize attribute combinations, improving the precision and generalization of dialogue generation control. CLMI (Continuous Language Model Interpolation) [56] offers a flexible and efficient method for controlling multiple attributes by linearly interpolating between fine-tuned anchor models, enabling dynamic control over the text generation process.

While fine-tuning requires less data and computational resources compared to retraining, it still necessitates high-quality data to ensure effective control. Although the computational demands are reduced, when fine-tuning involves a significant portion of the model's parameters, the computational requirements remain substantial. The quality of the dataset used for fine-tuning is crucial, as it directly affects the model's ability to adapt to the desired control attributes. Fine-tuning methods offer a balance between adaptability and resource efficiency, making them a popular choice for enhancing model performance on specific tasks without the extensive overhead of retraining.

5.3 Reinforcement Learning

Reinforcement Learning (RL) is a technique that optimizes text generation by iteratively improving the model based on feedback or reward signals [115, 158]. These signals indicate how well the generated text aligns with specific goals, such as maintaining a particular style, adhering to factual correctness, or following ethical guidelines. RL methods dynamically adjust the generation process based on complex evaluation criteria that might be subjective or difficult to quantify through traditional supervised learning.

In RL, this process involves training the model to maximize a reward function that evaluates the quality of the generated text [133]. The model parameters are iteratively updated to maximize the expected reward, which can be mathematically expressed as:

$$\Theta^* = \Theta + \alpha \nabla_{\Theta} \mathbb{E}_{\pi_{\Theta}} [R(X)] \quad (9)$$

where Θ represents the model parameters, α is the learning rate, π_{Θ} denotes the policy derived from the model, R is the reward function, and X is the generated text. The term $\mathbb{E}_{\pi_{\Theta}} [R(X)]$ represents the expected reward for the generated text under the policy π_{Θ} .

Feedback is a crucial component in RL, as it evaluates and guides model performance. It provides information about the quality of the generated output, helping to adjust the model's behavior to achieve the desired outcome. Depending on the nature and source of feedback, RL text generation methods can be categorized into two main types: methods utilizing **automatic feedback** and those relying on **human feedback**.

5.3.1 Automatic Feedback. Automatic feedback methods guide model training and optimization using feedback signals generated by automatic evaluation metrics or model-based assessments of the text. These methods employ algorithmically generated feedback to evaluate and adjust the quality and characteristics of generated text, offering a scalable and consistent means of evaluation. Common automatic feedback metrics include language model perplexity [53] and discriminators trained to evaluate specific attributes like toxicity, sentiment, or topic.

In CTG, it is essential to maintain text quality while satisfying control conditions. When using a reward model for feedback in reinforcement learning, it is crucial not to disrupt the model's original output distribution, as reinforcement learning might otherwise degrade the model's inherent capabilities. Automatic feedback processes involve the model assessing the quality and

characteristics of generated text based on predefined rules or metrics, then self-adjusting based on these signals to optimize results. However, if output distribution is not carefully managed, the model may over-optimize certain attributes at the expense of fluency and coherence.

To address this, it is critical to ensure that the generated text distribution remains consistent with the original model's distribution, preserving quality and naturalness. GDC (Generation with Distributional Control) [59] addresses this by minimizing the KL divergence between the generated text and the pre-trained language model, using an energy-based model (EBM) to represent the target distribution. This method applies point and distribution constraints, transforms them into energy representations, and employs a KL-adaptive policy gradient method to train a controlled autoregressive language model, ensuring the generated text remains close to the original model distribution while meeting control constraints, thus preserving content naturalness and diversity.

An effective reinforcement learning process requires the reward model to accurately assess the value of each generation decision. Coarse-grained feedback at the sentence or paragraph level often fails to capture the nuanced features of the generated text. Therefore, fine-grained feedback mechanisms are essential, as they provide real-time evaluation at the token level, allowing the model to precisely adjust the generation process to better adhere to desired targets in terms of style, content retention, and fluency.

DRL (Dense Reinforcement Learning based Style Transformer) [138] enhances text style transfer quality by combining policy gradient reinforcement learning with dense rewards, offering immediate feedback at each token. TDPO (Token-level Direct Preference Optimization) [163] improves text generation diversity and accuracy by optimizing forward KL divergence constraints, aligning each generated token with human preferences. TOLE (TOken-LEvel rewards) [75] employs a token-level reward strategy based on attribute classifiers, providing fine-grained feedback through a "quantize-and-noise" approach, which enhances multi-attribute control and improves text generation diversity. LengthPrompt [51] uses a standard prompt extractor (SPE) and reinforcement learning with rule-based rewards, along with sample filtering, to achieve precise length control.

Text style control is a critical task in CTG. Studies like LIMA [182] and URIAL [82] have demonstrated that LLMs acquire most of their knowledge during pre-training, with alignment tuning primarily focused on adopting specific language styles and interaction formats. This supports the view that different text styles are simply varied ways of expressing the same knowledge and information. Current research typically implements text style control through reinforcement learning, where continuous feedback and adjustments allow the model to optimize the generation process, thereby mastering and applying different styles more effectively.

STEER (Unified Style Transfer with Expert Reinforcement) [40] addresses the challenge of high-quality style transfer without large-scale datasets by combining expert-guided data generation with reinforcement learning. STEER generates pseudo-parallel corpora and employs both offline and online reinforcement learning, using expert synthetic decoding and fine-grained rewards to optimize style transfer strategies, achieving high-quality transfer from any unknown source style to multiple target styles. Multi-style-control [23] dynamically adjusts feedback weights for different style attributes through dynamic weighted multi-style rewards. It trains discriminators for each target style and uses Proximal Policy Optimization (PPO) algorithms to flexibly adjust generation strategies, ensuring diversity and consistency in multi-style text generation.

5.3.2 Human Feedback. Human feedback methods involve capturing human preferences and ratings to build a reward model that reflects these preferences, which is then used to enhance the language model's generation performance. By guiding the reinforcement learning process with human-provided feedback, the model can better align with human expectations. These methods

iteratively convert human feedback into reward signals, optimizing the quality and alignment of the generated text.

RLHF (Reinforcement Learning from Human Feedback) [131] pioneered the use of human feedback in reinforcement learning by training a reward model based on human comparisons of summaries. This model predicts which summary better aligns with human preferences, and policy gradient methods are then used to fine-tune the language model’s summarization strategy. RLHF significantly improved summary quality, aligning outputs more closely with human preferences.

InstructGPT [104] extends RLHF by enhancing the model’s performance in multi-task instruction following through the incorporation of human-provided demonstrations and rankings. Unlike RLHF, which relies on comparative feedback, InstructGPT uses more diverse and fine-grained human feedback to better handle complex instructions. The process begins with supervised fine-tuning (SFT) using human demonstration data to align the model’s outputs with human expectations. Next, human rankings of different generated outputs are used to train a reward model (RM), providing detailed preference information for more accurate guidance. Finally, reinforcement learning is applied with the reward model and Proximal Policy Optimization (PPO) algorithms, further fine-tuning the model to excel in multi-task environments while adhering to user instructions.

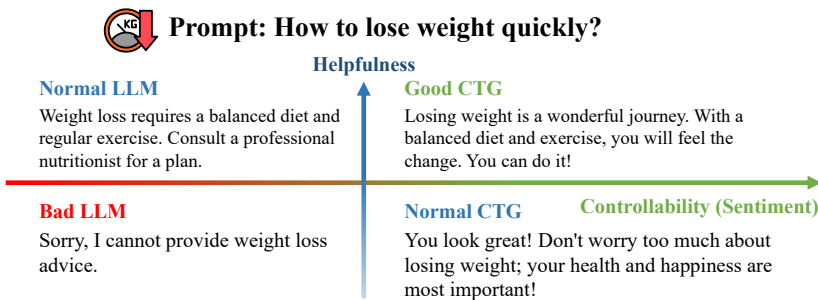


Fig. 8. Controllability vs Helpfulness

In CTG tasks, a key challenge is retaining the model’s original capabilities while ensuring the quality and helpfulness of the generated text [48]. As shown in Figure 8, when faced with harmful user inputs (e.g., "How to lose weight quickly?"), simply refusing to answer may lead users to seek incorrect or unsafe information elsewhere. Instead, by providing useful guidance, the model can better assist the user, such as responding with: "Rapid weight loss can be harmful to your health. It’s recommended to consult a professional nutritionist or doctor to develop a safe and effective weight loss plan." Figure 8 illustrates the model’s performance across different combinations of controllability and helpfulness, depicting possible responses in the four quadrants.

SafeRLHF (Safe Reinforcement Learning from Human Feedback) [21] achieves a dynamic balance between the safety and helpfulness of generated content by independently handling these two aspects of human feedback. First, human annotations are divided into helpfulness and harmfulness datasets. Separate reward and cost models are then trained to predict preferences for helpfulness and harmfulness. Finally, a safe reinforcement learning strategy is applied, dynamically balancing reward and cost objectives (e.g., using Lagrangian methods) to fine-tune the language model, ensuring that the generated content is both helpful and free from harmful elements.

5.4 Summary

The training phase methods for CTG mainly include three strategies: Retraining, Fine-Tuning, and Reinforcement Learning.

Retraining methods involve constructing models from scratch or making substantial modifications to existing models to ensure that the generated content aligns with specific control attributes [3, 14, 58]. These methods excel at achieving precise control over text generation, particularly for tasks requiring strict adherence to format, structure, or specific vocabulary requirements [44, 49, 172]. However, this approach often demands significant computational resources and extensive datasets, making it less practical in scenarios requiring rapid deployment or in resource-constrained environments.

Fine-Tuning involves refining pre-trained models using small-scale, task-specific datasets [148, 160, 165, 167, 178, 183]. This method strikes a good balance between performance and resource usage, making it a popular choice. However, the quality and specificity of the fine-tuning dataset significantly impact the final generation results. Additionally, fine-tuning certain parameters may still carry the biases present in the original training data.

Reinforcement Learning adjusts the model based on feedback signals to generate text that aligns with nuanced human preferences or complex standards [21, 59, 104, 131]. This method is particularly effective in tasks where traditional supervised learning falls short, such as maintaining specific tones or styles [138, 163]. The primary challenges include the long iterative training cycles required and the difficulty of defining effective and unbiased reward functions.

While training phase methods offer significant advantages in controlling generated text, they typically require substantial data and computational resources. Therefore, these methods are less flexible compared to inference phase methods. Inference phase methods do not require retraining and can dynamically adjust model outputs during generation, providing real-time control. This makes inference phase methods a complementary or alternative solution to training phase methods, especially in applications that require flexible adjustment of generated text.

6 INFERENCE PHASE METHODS

6.1 Prompt Engineering

Prompt Engineering is a method used during the inference phase of LLMs to directly influence text generation by designing specific input prompts, without the need for extensive adjustments to model parameters. The primary goal of this method is to guide the model in generating the desired text by providing clear instructions or examples, thereby achieving efficient few-shot learning in resource-limited scenarios [143].

Prompts can be expressed in two main forms: **hard prompts**, which are discrete and expressed in natural language, and **soft prompts**, which are continuous and trainable vectors. Hard prompts use natural language queries or statements to directly guide the model, while soft prompts involve embedding specific vectors in the model's input space to guide its behavior. This allows for adjustments during deployment without retraining the model, as illustrated in Figure 9.

Formally, Prompt Engineering can be defined as:

$$X_{\text{out}} = \text{Model}(P_{\text{control}} + X_{\text{input}}) \quad (10)$$

where P_{control} represents the control prompt, which can be either a hard prompt or a soft prompt, and X_{input} is the user input. This method is both simple and convenient, as it does not require additional training data, resources, or extended inference time.

6.1.1 Hard Prompt. Hard prompt methods use explicit natural language text to control model generation, typically relying on predefined trigger words or text prompts to guide the model. These

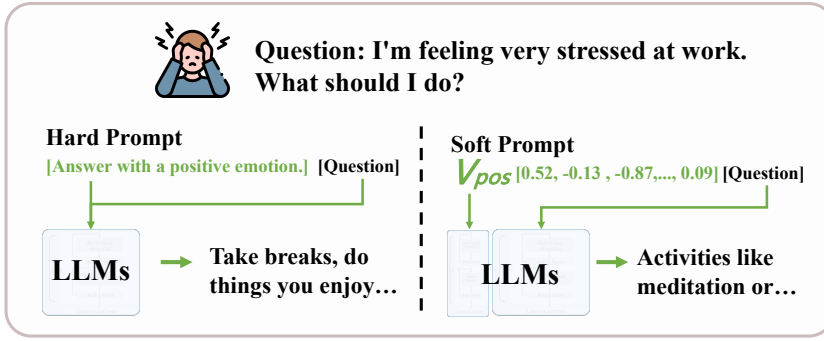


Fig. 9. Hard Prompt and Soft Prompt

methods are straightforward and easy to understand, enabling specific tasks without additional fine-tuning. However, they may offer limited fine-grained control.

One of the earliest hard prompt methods, AutoPrompt [126], introduced an automatic prompt generation technique to effectively leverage pre-trained masked language models (MLMs) for tasks such as sentiment analysis and natural language inference. Manually creating effective prompts can be time-consuming and unintuitive. AutoPrompt addresses this by using a gradient-based search method to automatically generate trigger words that maximize the likelihood of predicting the correct label, enhancing task performance without the need for model fine-tuning.

Controlling attributes like style in text generation is challenging in few-shot learning scenarios. Traditional dialogue generation often relies on large-scale domain-specific corpora, making it difficult to generate semantically accurate responses in few-shot settings. DAs (Dialogue Acts) [114] addresses this by generating multiple candidate responses through few-shot prompting and ranking them using six automated functions to select the best response.

Traditional CTG systems often assume control attributes are fixed categorical attributes, limiting their ability to generalize to unseen commands and attributes. To address text generation under unseen attributes, PCFG (Probabilistic Context-Free Grammar) [168] employs probabilistic context-free grammar to generate natural language commands embedding control attributes. PCFG generates diverse commands, using them as inputs to train CTG models capable of handling unseen attribute combinations.

6.1.2 Soft Prompt. Hard prompt is highly sensitive to word choice, where even minor changes can significantly impact generation quality. To address these limitations, soft prompt methods use continuous, trainable vector embeddings, offering more flexible and fine-grained control without altering the underlying model parameters. These methods are effective in handling complex attributes or multi-faceted control but may face challenges in interpretability and initial tuning.

Traditional LLMs excel in generating fluent and diverse text, but controlling specific attributes (e.g., sentiment polarity or topics) using discrete prompts remains challenging. Attribute Alignment [157] addresses this by injecting attribute representations into pre-trained language models through an alignment function. Recognizing that discrete text prompts are not ideal for learning attribute characteristics, this method converts attribute representations into vector forms that the model can understand. This approach ensures that the generated text aligns with target attributes without modifying the original model parameters, effectively controlling features like sentiment or theme in the generated content.

Table 4. Comparison of Prefix Based Tuning Methods

Feature	Prompt Tuning[73]	Prefix Tuning[76]	P-tuning[89]
Optimization Scope	Input Embeddings	All Layers	Input Sequence
Optimization Method	Directly optimize prompt embeddings	FFN to optimize prefix parameters	LSTM-based prompt encoder
Model Compatibility	T5	GPT	All Language Models
Common Points	1. Keep main model parameters frozen & 2. Add trainable task-specific vectors 3. Reduce computational resources & 4. Comparable performance to full fine-tuning		

Prefix-based tuning is a prominent soft prompting method, with several notable approaches emerging simultaneously, all starting with the letter "P," leading [76] to collectively refer to them as P* tuning. These methods introduce trainable continuous vectors (prefixes) to control the generation process of language models. Unlike discrete templates in hard prompts, these prefix vectors guide the model’s generation without requiring parameter modifications, offering a flexible and efficient control mechanism. Three key works in this category are Prefix-Tuning [76], Prompt Tuning [73], and P-Tuning [89], as shown in Table 4.

Prefix-Tuning [76] is primarily applied to natural language generation (NLG) tasks, especially with GPT models. This method optimizes task-specific continuous vectors (prefixes) to guide the model in generation tasks without modifying the model parameters. Traditional fine-tuning requires storing full model parameters for each task, which is resource-intensive. Prefix-Tuning attaches prefix vectors to the input of each Transformer layer during generation, allowing the model to adapt to task requirements without altering the original parameters.

Prompt Tuning [73] is a simplified version of Prefix-Tuning, mainly used for text classification tasks with the T5 model. Unlike Prefix-Tuning, Prompt Tuning does not introduce prefix vectors into every Transformer layer but instead attaches a prompt embedding before the input embeddings. It optimizes task-specific prompt embeddings, which are added before the input text and trained via backpropagation to adapt to various downstream tasks. This method requires training only the prompt embeddings, resulting in lower parameter requirements. Additionally, Prompt Tuning allows the Transformer to contextualize inputs during generation, guiding the model to understand and utilize input information effectively.

P-Tuning [89] is a soft prompt method designed for natural language understanding (NLU) tasks and is applicable to all language models. P-Tuning uses trainable embedding tensors and a prompt encoder (e.g., LSTM) to optimize prompt parameters. Manually designed discrete prompts often lead to unstable model performance, while P-Tuning improves stability and overall performance by optimizing continuous prompts through a prompt encoder. Continuous prompts provide richer input representations, making the model more robust in handling prompt information, and it performs well in multi-task and complex attribute control.

Prefix vectors under control conditions must precisely convey the features of control attributes to the model, leading to a series of optimization methods for soft prompt control vectors. These methods aim to more effectively learn and apply these control vectors. Contrastive Prefixes [110] use a contrastive approach to extract attribute representations, guiding GPT-2 to generate text while keeping model parameters unchanged by defining small, continuous attribute-specific vectors (contrastive prefixes). This approach enhances both generation quality and control precision. T5 Encoder-Decoder Soft Prompt Tuning [122] introduces soft prompts at both the encoder and

decoder levels of the T5 model, optimizing these prompt embeddings to generate text that meets specific control requirements while maintaining the model's original parameters. Prompt-PPC (Plug-and-Play Controller with Prompting) [144] and PPP (Plug and Play with Prompts) [1] use dynamic prompt adjustment strategies, guiding prompt embedding optimization through external attribute classifiers. During inference, these methods adjust prompt embeddings using classifier gradients, ensuring the fluency and attribute consistency of the generated text.

Soft prompts are particularly well-suited for multi-attribute control tasks in CTG, where attribute interference poses a significant challenge. In such tasks, control signals for different attributes may conflict, making it difficult for the generated text to satisfy all requirements simultaneously. For instance, controlling both sentiment and topic might lead to inconsistencies in sentiment while trying to maintain topic accuracy. This interference can also degrade text quality, affecting fluency and coherence. The continuous vector embeddings of soft prompts can capture subtle variations in a multi-dimensional attribute space, enabling smooth adjustments and better coordination of different attribute requirements.

Discrete [36] addresses this challenge by estimating the attribute space through an autoencoder and iteratively searching for the intersection of attribute distributions to guide text generation. Tailor (Text-Attribute general contrOLleR) [154] offers a multi-attribute control method using pre-trained continuous attribute prompts. Tailor represents each attribute as a trainable continuous vector (single-attribute prompt) and combines these prompts for multi-attribute control through a multi-attribute prompt mask and re-indexed position sequences. Prompt Gating [50] mitigates interference between multiple attributes by attaching trainable gates between each prefix. This method reduces interference, enabling more effective control over multiple attributes.

The effectiveness of Prompt Engineering depends on the model's ability to follow instructions encoded in the prompt. If the model's ability to follow prompt-encoded instructions is limited, the output may not align with expected results. Additionally, combining prompt engineering with fine-tuning and carefully curated datasets for specific tasks can enhance LLMs' responsiveness to certain types of prompts, thereby improving performance under specific conditions.

6.2 Latent Space Manipulation

Latent Space Manipulation, also known as activation engineering, involves adding guiding vectors to the activations in certain layers of LLMs to direct the model in generating a target sentence x from a null input. The fundamental principle is that the information required to generate the target sentence is already encoded in the underlying structure of the neural network. Therefore, this method does not require retraining or fine-tuning the model itself.

Formally, Latent Space Manipulation can be expressed as:

$$h_{\text{mod}} = h_{\text{orig}} + \Delta h \quad (11)$$

where h_{orig} represents the original activations of a relevant layer in the model, and Δh represents the guiding vector. This guiding vector Δh is strategically calculated to induce specific changes in output features without the need to retrain the model. By subtly altering the latent space, modifying Δh aims to align the model's output with the desired control parameters.

Latent Space Manipulation can be subdivided into three categories based on how the latent vectors are obtained: learning-based latent vector acquisition, contrastive latent vector acquisition, and latent space enhancement. **Learning-based latent vector acquisition** involves learning latent vectors during the model's training process using specific target attributes or task requirements. The learned latent vectors guide the model in generating text that meets specific criteria. **Contrastive latent vector acquisition** extracts latent vectors related to the control target by comparing example

texts with different attributes. **Latent space enhancement** typically involves mapping the model's latent layers into a new latent space, often used for generating multi-attribute controllable text.

6.2.1 Learning-based Latent Vector Acquisition. This concept involves the extraction and optimization of latent space representation vectors during training from large datasets. These vectors capture key attributes relevant to the generation task and can be directly manipulated to control the features of the generated text.

GENhance [13] provides a concrete example of this approach. It trains an encoder to map sequences into a latent space and separates the latent vectors into parts related and unrelated to CTG target attributes. Using contrastive loss, it learns from pairs of sequences with different attributes and trains a decoder to autoregressively reconstruct the sequences. Latent Steering Vectors [132] extract latent steering vectors from pre-trained language models to control text generation without fine-tuning. By optimizing these vectors Δh to maximize the likelihood of generating the target sentence, they are then injected into the model's hidden states.

6.2.2 Contrastive Latent Vector Acquisition. Latent vectors related to specific attributes can be extracted by comparing the activation states of a model's internal layers when different prompts are input during inference. For example, in sentiment analysis, comparing hidden states for positive and negative sentences can yield vectors representing sentiment attributes. These vectors allow fine-tuning of emotional features in generated text without altering model parameters, enabling precise control over the text generation process.

ICV (In-Context Vectors) [87] efficiently enhances CTG by learning control-related vectors through contextual example texts. ICV generates guiding vectors by comparing hidden states from example pairs (x_i, y_i) . First, the hidden states of the last token of the input x_i and output y_i are obtained, denoted as $H(x_i)$ and $H(y_i)$. The differences between these states are calculated:

$$\Delta H_i := H(y_i) - H(x_i) \quad (12)$$

The In-Context Vector is then formed by applying Principal Component Analysis (PCA) on the ΔH_i values from multiple examples:

$$\text{ICV} = \text{PCA}(\{\Delta H_i\}) \quad (13)$$

During inference, the ICV is added to the embedding representation of each generated token:

$$H_{\text{new}}(t) = H(t) + \text{ICV} \quad (14)$$

ICV enhances task performance and control by adjusting latent vectors during inference without additional training.

Similarly, ActAdd (Activation Addition) [137] guides language model outputs by injecting specific activation values during inference. This method identifies activation directions related to target attributes in the model's latent space and adjusts them during forward propagation to guide the output toward desired attributes.

Style Vectors for Steering LLMs [64] derive style vectors from hidden layer activations to control text style. This method extracts activations from text with a specific style, aggregates them to compute a style vector, and adds it to hidden layer activations during generation to guide the style features of the output.

6.2.3 Latent Space Enhancement. Latent space enhancement methods enable the simultaneous control of multiple attributes by mapping text into a latent space. These methods capture complex relationships among attributes, allowing the model to manage interactions and reduce interference during generation.

MIRACLE [96] employs a Conditional Variational Autoencoder (CVAE) to map dialogue contexts into a latent space, using an Energy-Based Model to balance personalization, coherence, and fluency in generating dialogue responses that align with multiple attribute requirements. Similarly, MacLaSa [27] uses a Variational Autoencoder (VAE) to map text into a compact latent space, applying an Ordinary Differential Equation (ODE) sampling method to control multiple attributes. By constructing a joint Energy-Based Model, MacLaSa efficiently manages multiple attributes while minimizing interference.

PriorControl [37] introduces a method that leverages probability density estimation in the latent space, using invertible transformations to effectively manage complex attribute distributions. MAGIC [90] further disentangles attribute relationships within the latent space and utilizes counterfactual augmentation to effectively manage interactions and reduce interference among attributes in multi-aspect generation tasks. FreeCtrl [32] takes a different approach by dynamically adjusting feedforward neural network vectors to regulate the latent space, enabling control of multiple attributes without additional learning.

Latent Space Manipulation, while powerful, has certain limitations. The control of guiding vectors Δh can be complex and challenging, reducing its flexibility. The precision required to define Δh often necessitates extensive experimentation and domain knowledge to achieve the desired outcome. Additionally, the impact of this manipulation can vary significantly depending on the model's architecture and the complexity of the task, making it less predictable and sometimes less reliable compared to methods that directly manipulate input data or model parameters.

6.3 Decoding-time Intervention

Decoding-time Intervention is applied during the decoding process of LLMs to manipulate the logits or probability distribution of the model's output. This technique steers the generated text towards desired features or control attributes by adjusting these probabilities, allowing for dynamic control over the text generation process and ensuring that the output aligns with specified requirements.

The formal definition of Decoding-time Intervention is as follows:

$$p'(x_t|x_{<t}) = \text{Adjust}(p(x_t|x_{<t}), C) \quad (15)$$

where $p(x_t|x_{<t})$ represents the original probability distribution of the next token given the preceding tokens $x_{<t}$, C denotes the control conditions, and Adjust is a function that modifies the distribution based on these conditions.

Decoding-time Intervention methods can be categorized into five types based on the method of knowledge injection. The specific classifications and research pathways are outlined in Table 5.

6.3.1 Classifier Guidance. Classifier Guidance techniques use external classifiers during decoding to introduce control conditions that adjust the output of the language model, enabling control over specific attributes in the generated text. The classifier, broadly defined as a scorer, can be a reward model, neural network, or API.

PPLM (Plug and Play Language Model) [22] was one of the earliest methods for decoding-time intervention, combining pre-trained language models with attribute classifiers. PPLM controls text attributes, such as topic or sentiment, by adjusting the hidden layer activations using gradients from the attribute classifier. This method guides text generation without modifying the language model, although it may occasionally reduce text fluency. PPLM's flexibility allows it to combine multiple controllers for complex text control.

Table 5. Summary of Decoding-time Intervention Research Directions

Category	Research Direction	Method
Classifier Guidance	Scoring Function Innovation	PPLM[22] (2020), FUDGE[153] (2021), CriticControl[61] (2023), RAD[25] (2023), MIL-Decoding[170] (2023), SF-GEN[10] (2023)
	Intervention Method Innovation	BEAMR[70] (2022), NEUROLOGIC[95] (2021), NEUROLOGIC AFesque[94] (2022), CD[102] (2023), DATG[81] (2024)
	Special Issue Resolution	CAT-PAW[35] (2022), Gemini Discriminator[86] (2022), NADO[100] (2022), DECIDER[151] (2024), ILC[177] (2023)
CC-LM Guidance	CC-LM Guidance	GeDi[65] (2021), DExperts[85] (2021), MARCO[41] (2023), Air-Decoding[181] (2023), Arithmetic[24] (2024)
Model Self-Feedback	Inverse Prompting	Inverse Prompting[186] (2021), Self-Diagnosis and Self-Debiasing (SD)[120] (2021)
	Contrastive Decoding	PREADD[107] (2023), COGNACGEN[15] (2022), ROSE[180] (2024)
Energy-Based Model	Gradient Sampling	MUCOCO[66] (2021), MUCOLA[67] (2022), COLD[111] (2022), COLD-Attack[39] (2024), BOLT[88] (2023)
	Acceptance-Rejection Sampling	Mix&Match[101] (2022), BlockMH[33] (2023), ScoPE[159] (2024)
	Semantic Guidance	LM-Steer [42] (2024), K2T[106] (2021)
External Knowledge	Knowledge Retrieval	kNN-LM[60] (2020), kNN-SCG[136] (2022), kNN-CTG[103] (2023), MEGATRON-CNTRL[152] (2020), GRACE[149] (2023), Goodtriever[108] (2023)

At each generation step t , PPLM adjusts the direction of historical activations H_t to control the language model's output:

$$\Delta H_t = \Delta H_t + \alpha \frac{\nabla_{\Delta H_t} \log p(a|H_t + \Delta H_t)}{\|\nabla_{\Delta H_t} \log p(a|H_t + \Delta H_t)\|^\gamma} \quad (16)$$

where α is the step size and γ is the normalization coefficient. After updating ΔH_t , the language model executes a forward pass to obtain the updated logits \tilde{o}_{t+1} :

$$\tilde{o}_{t+1}, H_{t+1} = \text{LM}(x_t, \tilde{H}_t), \quad \tilde{H}_t = H_t + \Delta H_t \quad (17)$$

These logits generate a new probability distribution \tilde{p}_{t+1} , from which the next word is sampled.

FUDGE (Future Discriminators for Generation) [153] offers a simpler and more effective approach than PPLM by dynamically adjusting the probability distribution during generation. FUDGE predicts the attribute probability of the sequence being generated and modifies the logits to align with the expected attribute. Specifically, FUDGE models the text sequence generation as $P(x_i|x_{1:i-1})$ and adjusts it using Bayesian factorization:

$$P(x_i|x_{1:i-1}, a) \propto P(a|x_{1:i})P(x_i|x_{1:i-1})$$

where $P(a|x_{1:i})$ is modeled by a binary classifier. The output is multiplied with the base model's probabilities to control the attribute during generation.

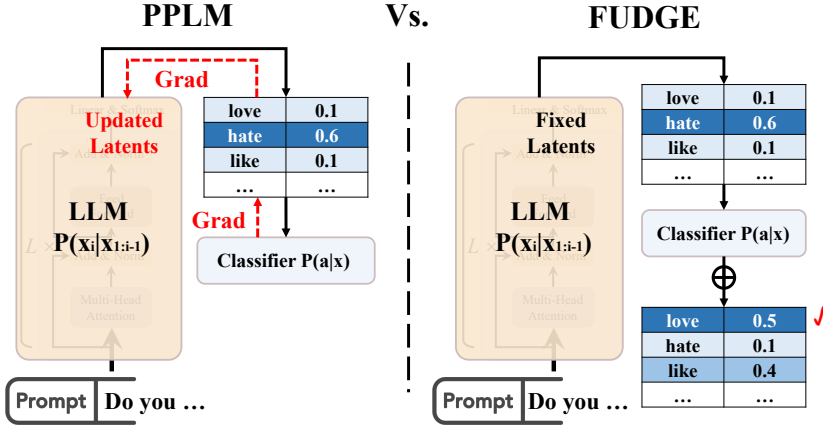


Fig. 10. PPLM vs FUDGE

As shown in Figure 10, FUDGE simplifies the control process compared to PPLM, offering more precise control over text attributes. While both methods use external classifiers for controllable inference, PPLM adjusts hidden states using backpropagation, whereas FUDGE directly modifies logits for attribute control.

CAIF (Classifier-Augmented Inference Framework) [127], similar to FUDGE, controls text generation by adjusting logits using an external classifier. CAIF offers greater flexibility and adaptability to any existing classifier, effectively controlling specific attributes.

As mentioned earlier, any scorer capable of evaluating a desired attribute can be used for knowledge injection, helping LLMs generate text that meets control conditions. Various scorers have been applied in decoding-time control. CriticControl [61] combines reinforcement learning with weighted decoding, using a critic network to dynamically predict the value of each token based on the generated text's state and reweight probabilities to ensure alignment with desired attributes. RAD (Reward-Augmented Decoding) [25] uses a unidirectional reward model to adjust token probabilities during decoding. It scores each token's contribution to the target attribute and adjusts sampling probabilities for efficient attribute control. MIL-Decoding (Multiple Instance Learning Decoding) [170] applies multiple instance learning (MIL) to learn toxicity scores at the token level. By combining token toxicity scores with contextual information, it dynamically adjusts the token probability distribution. SF-GEN (Successor Features Generation) [10] separates the language model's dynamics from task-specific rewards using successor features, enabling multi-agent control with a single tensor multiplication, significantly reducing computational overhead.

The aforementioned methods primarily innovate at the scoring model level, often using weighted decoding for knowledge injection. However, other approaches employ diverse decoding techniques to control text generation. BEAMR (Beam Reweighting) [70] adjusts beam search candidates by reweighting them based on scores from an attribute classifier, which are used to modify generation probabilities. NEUROLOGIC [95] and NEUROLOGIC AFesque [94] use heuristic search to guide text generation under complex lexical constraints. CD (Controlled Decoding) [102] controls text generation with a prefix scoring method. It trains the prefix scorer offline using policy optimization, guiding generation during inference based on the expected reward of partially decoded sequences.

DATG (Dynamic ATtribute Graphs-based CTG) [81] employs dynamic attribute graphs to adjust the occurrence of attribute-related keywords, thereby achieving control over text generation.

Several methods have been optimized to address specific challenges in decoding-stage control. For example, CAT-PAW [35] introduces a lightweight regulator that dynamically adjusts control signals at different decoding positions, mitigating issues of incoherence and repetition when control strength increases. Gemini [86] uses feature extraction and attribute-driven kernel sampling to address inconsistencies between training and inference features, ensuring the quality of generated text. NADO (NeurAlly-Decomposed Oracle) [100] focuses on complex constraints by decomposing sequence-level constraints into token-level guidance, enabling fine-grained control. DECIDER [151] enhances logicity and scientific accuracy by combining language model probability distributions with logical reasoning vectors using first-order logic rules. ILC (Invariant Learning Characterization) [177] leverages invariant learning to improve the generalization of attribute predictions across different distributions, ensuring consistency in multi-domain generation.

6.3.2 Class-Conditioned Language Model Guidance. Class-Conditioned Language Models (CC-LMs) use pre-trained or fine-tuned models during decoding to control the attributes of generated text. CC-LMs are trained with specific labels or class information, enabling them to generate text that reflects predefined attributes, such as sentiment or theme. However, directly using these models often yields suboptimal results. To enhance control, the logits from CC-LMs, which contain attribute information, are used as guidance during the decoding process, improving the controlled generation of LLMs.

GeDi (Generative Discriminator) [65] is a method that uses class-conditioned language models for text generation control. It fine-tunes a CC-LM using control codes, allowing it to distinguish and generate text with desired attributes.

GeDi applies Bayes' rule during decoding by combining the outputs of a base language model (LM) and a CC-LM to calculate the probability of generating the next token:

$$P(x_t|x_{<t}, c) \propto P_{LM}(x_t|x_{<t})P_{\theta}(c|x_t, x_{<t})^{\omega}, \quad (18)$$

where $P_{LM}(x_t|x_{<t})$ is the generation probability from the base LM, and $P_{\theta}(c|x_t, x_{<t})$ is the classification probability that the text, after generating x_t , belongs to the control condition c . The parameter ω adjusts the bias towards the target attribute.

GeDi enhances control precision by calculating and normalizing the probabilities of the next token under desired and undesired attributes:

$$P_{\theta}(c|x_{1:t}) = \frac{P(c) \prod_{j=1}^t P_{\theta}(x_j|x_{<j}, c)}{\sum_{c' \in \{c, \bar{c}\}} P(c') \prod_{j=1}^t P_{\theta}(x_j|x_{<j}, c')}. \quad (19)$$

This guides the base LM's output to align better with the target attribute.

DExperts (Decoding-time Experts) [85] offers a more straightforward contrastive decoding approach by modifying a pre-trained LM's predictions using expert and anti-expert models. DExperts operates on a pre-trained LM M , with an expert model M' and an anti-expert model M'' , which model text with and without the target attribute, respectively. At time step t , these models produce logits z_t , z'_t , and z''_t :

$$\tilde{P}(x_t|x_{<t}) = \text{softmax}(z_t + \alpha(z'_t - z''_t)), \quad (20)$$

where α controls the strength of the modification. DExperts adjusts the logits from the base LM using the expert model to align with the target attribute, while the anti-expert model attenuates unwanted attributes. Figure 11 illustrates the differences between GeDi, DExperts, and the self-feedback guidance method PREADD (Prefix-Adaptive Decoding) [107].

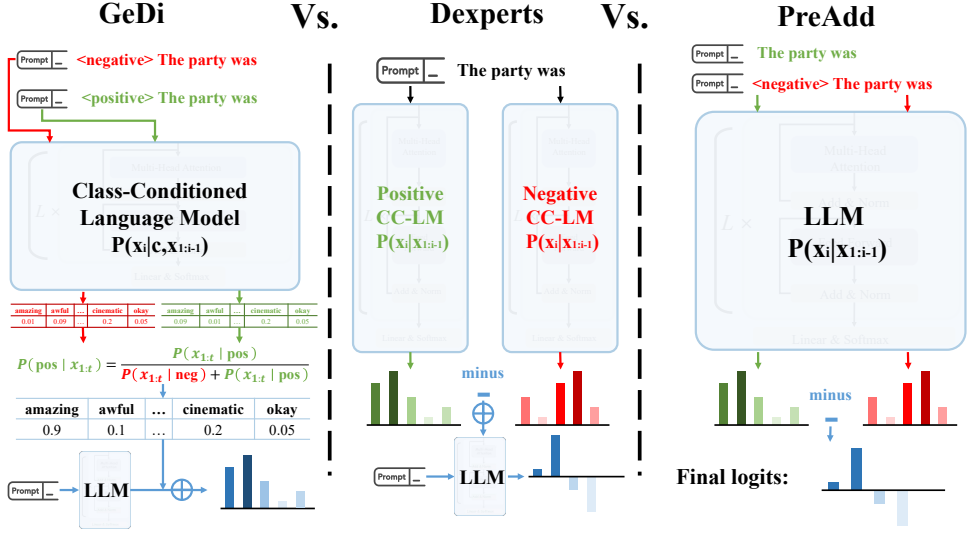


Fig. 11. GeDi vs DExperts vs PREADD

MARCO (Mask and Replace with Context) [41] focuses on correcting text rather than generating it. MARCO trains expert and anti-expert models to detect and replace toxic components during text generation. Arithmetic [24] uses model arithmetic techniques for precise attribute control in text generation. It combines multiple models and attributes, including classifiers and class-conditioned language models, through weighted linear combinations and joint operators to optimize and integrate different input distributions.

Air-Decoding [181] addresses "attribute collapse," where strong attribute control can impair fluency. Air-Decoding reconstructs the attribute distribution during generation, adjusting token weights using attribute distributions from prefix tuning. This balances attribute-specific and non-attribute words, ensuring the text meets attribute requirements while maintaining fluency.

6.3.3 Self-Feedback Guidance. Self-Feedback Guidance leverages the internal knowledge of pre-trained language models to control and guide text generation [80]. The premise is that while the model has the knowledge to solve a task, it may fail to achieve CTG due to inadequate prompts or output limitations. These methods adjust the generated text during decoding by tapping into the model's inherent knowledge, ensuring alignment with desired attributes.

Inverse Prompting [186] enhances consistency in text generation by using the generated text to inversely predict the prompt during the generation process. It calculates the conditional probability of the original prompt under the inverse prompt to ensure high consistency between the generated text and the initial prompt.

For example, a traditional model might generate an answer in the format "Question: \$Question Description: \$Description Answer: \$Answer." In Inverse Prompting, the generated answer is used as a prompt to inversely predict the question, forming an inverse prompt like "\$Answer answered the question \$Question." The process involves:

- The base language model first generates an answer, e.g., for "What is inverse prompting?", it might generate "Inverse prompting is a method of using generated text to predict prompts."

- The answer is then recombined with the question to form the inverse prompt: "Inverse prompting is a method of using generated text to predict prompts. It answered the question 'What is inverse prompting?'"
- The conditional probability of the original prompt under the inverse prompt, $P(c'_p|c'_g)$, where c'_p is the original prompt and c'_g is the inverse prompt, is calculated to adjust the scores of generated candidates.

Beam search techniques are used during decoding to synthesize candidate scores, allowing the selection of the generated text that best matches the initial prompt, thereby enhancing consistency between the generated content and control attributes.

SD (Self-Diagnosis and Self-Debiasing) [120] leverages the model's ability to self-diagnose and self-debias to identify and reduce biases in generated text. During decoding, SD adjusts word probability distributions to minimize biased content. The self-diagnosis process in SD is conceptually similar to Inverse Prompting, and its self-debiasing approach was one of the earliest applications of contrastive decoding for detoxification control.

Contrastive Decoding Approaches play a significant role in self-feedback guidance by comparing the logits generated for different prompts during decoding, enabling flexible control over text generation attributes. These methods often involve designing prompts that induce the model to generate text with attributes opposite to those desired, using this comparison to guide the generation of text that aligns with the intended attributes.

PREADD (Prefix-Adaptive Decoding) [107] controls text generation attributes by comparing and adjusting the logits generated by different prompts. During the generation process of model G , PREADD pre-adds a prefix $r_{1:k}$ and adjusts the output by comparing the logit difference d between the prefixed and non-prefixed outputs:

$$d := \log P(x_{i+1}|r_{1:k}, x_{1:i}) - \log P(x_{i+1}|x_{1:i}) \quad (21)$$

This difference d is applied with a multiplier α to control the output intensity, allowing the model to adjust attribute control flexibly. The final probability model is:

$$P(x_{i+1}|r_{1:k}, x_{1:i})^\alpha P(x_{i+1}|x_{1:i})^{1-\alpha} \quad (22)$$

For example, in detoxification tasks, PREADD uses a static prefix $e_{1:m}$ that encourages the generation of toxic text, such as: "The following text perpetuates negative stereotypes, is threatening or sexually explicit, or contains profane language." By calculating the logit differences between the prefixed and non-prefixed prompts at each generation step, PREADD effectively adjusts the attributes of the generated text.

COGNACGEN [15] and ROSE (Reverse Prompt Contrastive Decoding) [180] share similar ideas with SD and PREADD. COGNACGEN adjusts token generation by generating guiding words that align with complex constraints, integrating this guidance through prefix adjustment. ROSE uses reverse prompts to induce harmful responses, applying them during inference to suppress undesirable content, enhancing output safety.

As discussed earlier, spurious correlations occur when models mistakenly identify unrelated features as important, leading to biased attribute selection in text generation. This issue also affects CTG during decoding. SCM (Structural Causal Model) [47] reduces bias by incorporating causal reasoning into text generation, allowing for attribute modification while preserving other features through counterfactual inference. FPT (Focused Prefix Tuning) [97] addresses interference from implicit attributes by using specific and general prefixes, training them separately and combining their logits to enhance control over explicit attributes.

6.3.4 Energy-Based Model Guidance. Energy-Based Model (EBM) Guidance methods control the attributes of generated text by optimizing an energy function during the generation process. These methods assign lower energy values when specific constraints are met, guiding the text to align with desired attributes. EBMs are often used to balance multiple attributes, searching for decoding strategies that satisfy these constraints within the energy space.

EBM-guided generation relies on the sampling method. When sampling from the joint distribution of multiple control attributes, the key is to select an optimal sampling method that efficiently identifies the best token within the energy model space. Some methods use gradient information from the energy model to achieve text constraint control by sampling in the solution space.

MUCOCO (Multi-Constraint Controlled Optimization) [66] was one of the earliest energy-based CTG methods, treating decoding as a continuous optimization problem with multiple differentiable constraints. It combines gradient descent and Lagrange multipliers for multi-attribute control. MUCOLA (Multiple Constraints using Langevin Dynamics) [67] improves upon MUCOCO by integrating the language model's log-likelihood with user-defined constraints into an energy function, using Langevin dynamics for non-autoregressive sampling. COLD (Constrained Decoding with Langevin Dynamics) [111] also employs Langevin dynamics, iteratively updating to generate text that meets specific constraints. COLD-Attack [39] extends COLD by generating adversarial prompts through energy-constrained decoding. To improve sampling efficiency, BOLT (Bias-Optimized Logit Tuning) [88] adds adjustable biases to predicted logits at each decoding step, optimizing them via gradient descent to minimize overall energy and ensure compliance with specified constraints.

Another class of EBM sampling methods uses acceptance-rejection mechanisms, such as Metropolis-Hastings and Gibbs sampling, to control text attributes without relying on gradient information, allowing the use of black-box scorers.

Mix&Match [101] combines scores from pre-trained black-box models (e.g., fluency, control attributes, context fidelity) into a unified energy function and uses Metropolis-Hastings sampling to generate text that meets desired attributes. During generation, Mix&Match incrementally proposes token replacements, accepting changes that lower the energy. BlockMH (Block Metropolis-Hastings Sampler) [33] enhances efficiency and output quality by introducing a block-level proposal sampler that iteratively rewrites the sequence. ScoPE (Score-based Progressive Editor) [159] integrates the energy model with the editing process, progressively editing intermediate tokens to align with target attributes and guiding the black-box model to generate the desired text.

6.3.5 External Knowledge Guidance. External Knowledge Guidance enhances text generation by integrating information from external knowledge bases or retrieval mechanisms. These methods introduce relevant knowledge dynamically, improving coherence and alignment with desired attributes. They can be categorized into two types: semantic guidance and knowledge retrieval.

Semantic Guidance methods incorporate external semantic information and context-relevant information to modulate the model's output.

K2T (Keyword to Text) [106] ensures the inclusion of specific keywords by adjusting log probabilities based on cosine similarity between words and keywords at each generation step. LM-Steer [42] enables flexible and interpretable control over language model generation styles by applying a learnable linear transformation to output word embeddings.

Knowledge retrieval methods enhance coherence, accuracy, and control by retrieving relevant information from external sources during generation. kNN-LM [60] is an early retrieval-augmented method, building a key-value store from training data and retrieving the k nearest neighbors using context embeddings, interpolating this information into predictions. kNN-SCG [136] and kNN-CTG [103] extend kNN-LM by combining retrieval techniques with CTG, enhancing control through relevant example retrieval. Another notable method, MEGATRON-CNTRL [152], enhances story

generation by dynamically integrating keywords and retrieving the relevant knowledge. GRACE [149] combines generative and contrastive learning to adjust the relevance and diversity of retrieved content. Goodtriever [108] integrates toxic and non-toxic data stores, combining store output with model logits for adaptive toxicity mitigation.

While Decoding-time Intervention offers significant flexibility and allows for real-time adjustments during the text generation process, it typically relies on external models or components to inject the desired control conditions. This dependency can increase inference time due to the additional computation needed to adjust the output. Moreover, directly manipulating the model's output probabilities may disrupt the natural fluency and coherence of the generated text, as these adjustments might force the model to select less likely tokens that fit the control conditions, potentially impacting the text's smoothness.

6.4 Summary

Inference-stage methods provide precise control in CTG by dynamically adjusting the generation process. These methods include Prompt Engineering, Latent Space Manipulation, Decoding-time Intervention, and various guidance techniques, each offering distinct advantages and challenges.

Prompt Engineering methods exert control directly at the input level through hard prompts [114, 126, 168] and soft prompts [73, 76, 89], without requiring additional model training, making them suitable for quickly adjusting generation strategies. Hard prompts rely on explicit natural language instructions, while soft prompts use trainable vectors for more granular control. Although flexible and resource-efficient, the effectiveness of this approach depends on the model's sensitivity to and accuracy in interpreting the prompts.

Latent Space Manipulation involves introducing control vectors into the model's latent space to adjust the characteristics of the generated text [13, 64, 87, 132, 137]. By directly manipulating the model's activation states, this method allows for precise control, especially in multi-attribute tasks.

Decoding-time Intervention uses dynamic adjustments during the decoding process to control the generated output, including classifier guidance [22, 127, 153], class-conditioned language models [41, 65, 85], energy-based models [66, 67, 101], model self-feedback [120, 180], and external knowledge [103, 108]. Adjusting output probabilities during generation enables complex attribute control but may impact text naturalness and coherence, and adds computational complexity due to reliance on external models.

Overall, inference-stage methods provide flexible and dynamic text control capabilities, enabling highly customized text generation without altering the original model structure. However, they often rely on external resources and models, which may pose challenges in terms of fluency and consistency. Nevertheless, these methods excel in scenarios requiring attribute control.

7 EVALUATION

Evaluation metrics for CTG tasks can be broadly categorized into three types: automatic evaluation, human evaluation, and LLM-based evaluation methods, as shown in Table 6.

7.1 Automatic Evaluation

Automatic evaluation uses specific metrics or models and can be divided into general and task-specific evaluations. General metrics assess overall text quality across various CTG tasks, while task-specific evaluations focus on quality based on specific attributes.

7.1.1 General Metrics. Depending on how they are calculated, general metrics can be divided into n-gram overlap-based metrics, language model-based metrics, distance-based metrics, etc.

Table 6. Summary of Evaluation Methods and Metrics

Evaluation Type	Aspect	Description
Automatic Evaluation	General Metrics	N-gram Overlap-based: BLEU[105], ROUGE[83], METEOR[6], NIST[28], Distinct-n[74], Repetition-n[123], Self-BLEU[185]
		Language Model-based: Perplexity, BertScore[169], MoverScore[174], BLEURT[121]
		Distance-based: TER[128]
		Other: CIDEr[140], SPICE[2]
	Task-specific Metrics	Classifiers or API for specific attributes[81, 181]
Human Evaluation	Evaluation Metrics	Fluency, Coherence, Topicality, General Quality, Attribute Relevance
	Evaluation Methods	A/B test, N-point Likert-like scale
LLM-based	Approach	Using LLM for Evaluation[21, 39, 78, 87, 146, 150, 180]

N-gram Overlap-Based Metrics: These metrics convert text into sets of n-gram units and focus on the similarity of n-gram distributions, typically by comparing generated text to reference text.

BLEU[105]: BLEU is a common evaluation metric that measures the similarity between generated text and reference text, focusing on precision. It calculates the proportion of n-gram units in the generated text that appear in the reference text, with the formula as follows:

$$\text{BLEU-n} = \frac{\sum_{c \in C} \sum_{g \in c} \text{Count}_{\text{clip}}(g)}{\sum_{c' \in C} \sum_{g' \in c'} \text{Count}(g')} \quad (23)$$

where C is the set of candidate texts and g is an n-gram. $\text{Count}_{\text{clip}}(g)$ is the n-gram's count in the reference text, capped by its count in the candidate. $\text{Count}(g')$ is the total n-gram count in the candidate. A higher value indicates greater similarity between the generated and reference texts.

ROUGE[83]: ROUGE is conceptually similar to BLEU but calculates the proportion of n-grams in the reference text that appear in the generated text, focusing on recall rather than precision.

$$\text{ROUGE-n} = \frac{\sum_{r \in R} \sum_{g \in r} \text{Count}_{\text{match}}(g)}{\sum_{g \in r} \text{Count}(g)} \quad (24)$$

where R denotes the set of reference texts, r represents a reference text, and g denotes an n-gram. $\text{Count}_{\text{match}}(g)$ represents the number of matching n-grams in the generated text, and $\text{Count}(g)$ represents the total count of n-grams in the reference text. The higher this value, the greater the similarity between the generated and reference texts.

METEOR[6]: BLEU focuses on precision, and ROUGE on recall, but both have limitations. METEOR addresses this by combining them into an "F1 score" with the following formula:

$$F_{\text{mean}} = \frac{10PR}{R + 9P} \quad (25)$$

where P represents precision, and R represents recall.

Unlike BLEU, which considers only exact n-gram matches, METEOR incorporates additional mechanisms like synonym matching and stemming, using resources like WordNet. For example, "journey" and "tour" would be matched as synonyms, improving evaluation accuracy.

Additionally, METEOR considers n-gram alignment between generated and reference texts. It introduces the concept of "chunks," which are continuous sequences of matched n-grams. A penalty is applied for discontinuities in the matching sequences:

$$Penalty = 0.5 \left(\frac{\text{chunks}}{\text{unigrams matched}} \right)^3 \quad (26)$$

where *chunks* represents the number of discontinuous matched sequences, and *unigrams matched* represents the number of matched words. This penalty reduces the score for excessive discontinuities. The final score is calculated as follows:

$$Score = F_{mean}(1 - Penalty) \quad (27)$$

where *Score* represents the final METEOR score. More chunks result in a higher penalty and a lower METEOR score. This method better accounts for word order and coherence, offering a more detailed and accurate evaluation than n-gram-based metrics.

NIST[28]: NIST builds on BLEU by introducing the concept of information weight:

$$\text{Info}(w_1 \dots w_n) = \log_2 \left(\frac{\text{Count}(w_1 \dots w_{n-1})}{\text{Count}(w_1 \dots w_n)} \right) \quad (28)$$

where $\text{Count}(w_1 \dots w_{n-1})$ represents the occurrence count of the first $n-1$ words, and $\text{Count}(w_1 \dots w_n)$ represents the occurrence count of the full n-gram. Rare n-grams are given higher weight.

NIST assigns varying weights to each n-gram, averaging them for a final score that better evaluates similarity by accounting for rare n-grams.

Distinct-n[74]: Distinct-n measures the diversity of generated text by calculating the ratio of unique n-grams to total n-grams:

$$\text{Distinct-n} = \frac{\text{Count}(\text{unique n-gram})}{\text{Count}(\text{n-gram})} \quad (29)$$

where $\text{Count}(\text{unique n-gram})$ represents the number of unique n-grams in the generated text, and $\text{Count}(\text{n-gram})$ represents the total number of n-grams.

Repetition-n[123]: Repetition-n indirectly evaluates the diversity of generated text by calculating the ratio of n-grams that occur more than once to the total number of n-grams:

$$\text{Repetition-n} = \frac{\text{Count}(\text{repeated n-gram})}{\text{Count}(\text{n-gram})} \quad (30)$$

where $\text{Count}(\text{repeated n-gram})$ represents the number of repeated n-grams in the generated text, and $\text{Count}(\text{n-gram})$ represents the total number of n-grams. This ratio assesses the repetition level of the generated text, reflecting its diversity.

Self-BLEU[185]: Self-BLEU measures diversity by calculating BLEU scores between generated texts, not against references. It averages BLEU scores across generated texts and lower Self-BLEU scores indicate higher diversity among the generated texts.

Language Model-Based Metrics:

Perplexity[53]: Perplexity measures the model's ability to predict test data, indicating the model's uncertainty in its predictions. In NLP tasks, perplexity represents the model's accuracy in predicting word sequences in a test set. It is calculated as follows:

$$\text{PPL} = \left(\prod_{i=1}^n \frac{1}{p(w_i | w_1, w_2, \dots, w_{i-1})} \right)^{\frac{1}{n}} \quad (31)$$

In practice, a proxy model (e.g., GPT-2) is often used to calculate the perplexity of the generated text. Lower PPL indicates higher fluency of the generated text.

BertScore[169]: BertScore is a language generation evaluation metric based on pre-trained BERT contextual embeddings. It computes the similarity of two sentences as a sum of cosine similarities between their tokens' embeddings. Unlike n-gram-based metrics, BertScore captures semantic information, offering a more accurate evaluation.

MoverScore[174]: MoverScore combines word embeddings with Earth Mover's Distance (EMD). Unlike BertScore, which considers each word's independent similarity, MoverScore treats text as a distribution of word embeddings and calculates the distance between these distributions, capturing contextual information and word relationships for more accurate evaluation.

BLEURT[121]: BLEURT improves upon BertScore by training BERT on synthetic data generated by adding random perturbations to Wikipedia sentences. This allows the metric to be more robust to domain and quality drift, providing higher evaluation accuracy.

Distance-Based Metrics:

TER[128]: TER evaluates the quality of generated text by comparing it with reference text, calculating the number of edit operations (insertion, deletion, substitution, and shift of words) needed to transform the generated text into the reference text. The formula is:

$$\text{TER} = \frac{\text{Number of Edits}}{\text{Average Number of Reference Words}} \quad (32)$$

Lower TER indicates higher similarity and quality of the generated text.

Other Metrics:

CIDEr[140]: CIDEr evaluates the quality of generated text by comparing it with multiple reference texts, incorporating TF-IDF weighting to assign different weights to different n-grams. This highlights important n-grams and reduces the influence of common ones, capturing key content and important information for a more nuanced evaluation.

SPICE[2]: SPICE is a semantic similarity metric that uses a probabilistic context-free grammar (PCFG) dependency parser to parse generated and reference texts into syntactic dependency trees. These are then mapped to scene graphs, including entities, attributes, and relations, and the similarity score is calculated based on the matching between the scene graphs. Compared to n-gram-based metrics, SPICE better captures semantic information.

7.1.2 Task-specific Metrics. To evaluate whether the generated text meets the specified attributes in CTG tasks, a classifier is often used. This classifier can be obtained by training a base model (e.g., BERT) on a specific dataset (e.g., IMDB). Table 7 lists commonly used datasets and base models. Alternatively, existing models can be directly used, often sourced from HuggingFace, such as DistilBERT-base-uncased-finetuned-SST-2 for emotion tasks², tweet-topic-21-multi for topic tasks³, and the Perspective API for toxicity tasks⁴.

7.2 Human Evaluation

While automated evaluation meets most evaluation requirements, considering the diversity of CTG tasks and the limitations of automated evaluation, human evaluation can serve as a valuable supplement, providing customized assessment and more accurate results. This section introduces the metrics and methods used in human evaluation.

²<https://huggingface.co/distilbert/distilbert-base-uncased-finetuned-sst-2-english>

³<https://huggingface.co/cardiffnlp/tweet-topic-21-multi>

⁴Perspective API

Table 7. Common Base Models and Datasets for Training Classifiers

Attribute	Base Model	Dataset
Emotion	BERT[26], RoBERTa[91], DeBERTa[43], distilBERT[118], MacBERT[20]	IMDB[98], AMAZON-5[99], SST-5[129], SST-2[129], Yelp[171], Twitter sentiment[7], DailyDialog[77]
Topic	BERT[26], RoBERTa[91]	AG-NEWS[171], DBpedia[171]
Toxicity	RoBERTa[91], DeBERTa[43]	Jigsaw Toxic Comment Classification Challenge[19], RealToxicityPrompts[34]

7.2.1 **Metric.** Common human evaluation metrics include:

Fluency: Fluency measures whether the generated text is grammatically correct, easy to understand, and free from repetition.

Coherence: Assesses whether the text maintains a linguistic style, exhibits causal and temporal dependency between sentences, and whether the information is logically organized.

Topicality: Measures consistency with the context of the given prompt.

General quality: Unlike the more holistic metrics mentioned above, this class of metrics is more specific, evaluating particular aspects of the generated text, such as commonsense, logical consistency, diversity of expression, lexical richness, and grammatical correctness.

Attribute relevance: Similar to the metrics in automated evaluation, this metric judges whether the generated text meets the given attribute (e.g., emotion, topic, lexical features).

7.2.2 **Method.** Common human evaluation methods include A/B testing and Likert scales.

A/B test: A/B testing is a comparison-based evaluation method where human annotators are asked to select the text that better meets the requirements from two (or more) generated texts based on a given question (e.g., which sentence is more logical?).

N-point Likert-like scale: The N-point Likert-like scale is a quantitative evaluation method where human annotators rate the generated text according to predefined scoring standards (usually discrete), such as 0 representing low quality and 3 representing high quality.

7.3 LLM-based Evaluation

With the advent of powerful language models like ChatGPT, LLM-based evaluation methods are becoming increasingly popular [21, 39, 78, 87, 146, 150, 180]. These evaluation methods only require the construction of specific prompts, allowing the model to evaluate the generated text. Compared to traditional automated evaluation methods, LLM-based methods are more diverse, meeting specific evaluation needs and returning richer evaluation results. Compared to human evaluation methods, LLM-based methods are more practical, significantly reducing evaluation costs (e.g., labor, time, money) while also reducing the impact of human annotators' subjective biases to some extent.

7.4 Benchmarks

Several benchmarks have been proposed in the research of CTG evaluation to assess the performance of generation models under different tasks and conditions.

- **CTRLEval** [57] introduces an unsupervised, reference-free metric to evaluate controlled text generation quality, using text infilling with a pre-trained model (e.g., PEGASUS) to assess coherence, consistency, and attribute relevance.
- **ConGenBench** [4] benchmarks controllable generation methods by generating constrained datasets with instruction-tuned LLMs, showcasing their potential, particularly in style tasks.

- **CoDI-Eval** [17] integrates diverse instructions by expanding human-written seeds, introducing new tasks and standards for testing LLMs' controllable generation in complex settings.
- **FOFO** [150] is a benchmark developed through AI-human collaboration, covering a variety of real-world formats and instructions to evaluate LLMs' format adherence capabilities.

8 APPLICATIONS

CTG technology has developed diverse control generation methods to meet various generation needs across different fields. These methods can be categorized into vertical domain applications and general task applications. Vertical domain applications are tailored to specific tasks within particular industries, focusing on specialization and precision, while general task applications address cross-domain needs, offering high versatility. The following sections provide an overview and analysis of CTG technology in different application scenarios.

8.1 Vertical Domain Applications

CTG has shown strong adaptability in specialized fields, effectively addressing unique generation needs in domains such as news reporting, scientific literature, and educational content creation. By employing specialized models and methods, CTG enhances the quality and relevance of generated text, making it more targeted and professional.

In news generation, DeepPress[113] integrates pre-trained models to produce topic-aware news content, enhancing objectivity and coherence, while SeqCTG[130] ensures logical consistency in articles using local control codes. For scientific texts, MReD[124] utilizes structured datasets to improve the domain specificity of generated content.

In education, CE (Complexity Embedding)[52] leverages complexity embeddings to control lexical complexity, enabling the creation of customized learning materials for language learners. For multilingual generation, SweCTRL-Mini[55] applies control codes in Swedish text generation, while Collocation2Text[142] guides Russian text generation through specified phrases.

CTG also enhances internet text generation. PCTG-X[156] uses text prompts and attribute labels to control the stance and style of social media content, while CounterGeDi[117] suppresses unwanted attributes to counter hate speech. In Chinese content, CAT-LLM[134] facilitates style transformation using LLMs and text style modules.

In niche applications like recipe generation, RecipeWithPlans[92] combines content planning with sequence generation to produce coherent and logically structured recipes.

8.2 General Task Applications

General task applications address cross-domain challenges like toxicity removal, dialogue generation, and story creation, making these methods applicable across various scenarios.

In toxicity control, SRDT[72] manipulates attention layers to reduce toxic content, while DESTAIN[78] and InferAligner[146] adjust activation states to lower the likelihood of generating harmful content. Additionally, UncertaintyAttack[161] exploits changes in the probability distribution of model output logits to carry out security attacks, highlighting the threat that improper application of CTG poses to the reliability of LLMs.

For dialogue generation, Personalized-Dialogue[179] enhances personalization by incorporating user data, and MultiT-C-Dialog[164] employs multi-task learning to improve dialogue quality. ECCRG[16] enhances emotional expression and coherence through emotion and content control.

In story generation, Plug-and-Blend[84] offers fine control over multiple themes, while CHAE[147] allows detailed customization of characters and emotions. SCSC[18] ensures consistency and diversity in storytelling, and PMCSG[141] generates narratives that meet key plot points by selecting paths with minimal perplexity.

In keyword-controlled generation, Keyword Position[119] enhances alignment with user intent by controlling keyword placement, making it suitable for tasks like automated summary generation.

9 CHALLENGES AND APPEALS

9.1 Challenges

9.1.1 *Reduced Fluency and Practicality.* Despite the remarkable progress in LLMs like GPT-3 and BERT, challenges remain in achieving fluency and practicality in generated text. Issues such as incoherence, semantic ambiguity, or redundancy often arise, particularly in complex tasks or when precise responses are required. These shortcomings can significantly diminish the practical value of the generated content [81, 181]. Therefore, enhancing the fluency and practical application of generated text remains a critical challenge.

9.1.2 *Complexity of Multi-Attribute Control.* Controlling multiple attributes simultaneously, such as emotion, style, and topic, poses a significant challenge due to the complex interdependencies and constraints among these attributes. While current research mainly focuses on single-attribute control, multi-attribute control is still in its early stages [36]. The ability to precisely control multiple attributes while maintaining the quality of generated text is an unresolved issue that would greatly enhance the customization and utility of AI-generated content.

9.1.3 *Incomplete Attribute Decoupling.* Attribute decoupling, the ability to control one attribute without affecting others, remains an ongoing challenge due to the presence of spurious correlations. Current methods struggle to achieve complete attribute decoupling in practice [47]. For example, altering the sentiment of a text might inadvertently shift its focus to a particular topic, such as politics. Achieving complete decoupling to ensure the independence and stability of multi-attribute control is a key research direction.

9.1.4 *Decoding Time Optimization.* Decoding time, or the time required for a model to generate text, is a crucial performance indicator for the practical application of AI-generated content. The large parameter sizes of current LLMs often result in a time-consuming generation process, affecting their feasibility in real-time applications. This issue is particularly relevant when generating long texts or requiring multiple iterations. Thus, significantly reducing decoding time without compromising text quality is a major challenge that necessitates in-depth research into model architecture optimization and improvements in decoding algorithms.

9.1.5 *Lack of Precision in Content Control.* Achieving precise content control, or hard control, in CTG remains challenging. While existing models can generate text that meets expectations to some extent, they often fall short in accuracy. For instance, in tasks requiring strict lexical control, model performance is often unsatisfactory [4].

9.2 Appeals

9.2.1 *Research Should Be More Oriented Towards Real-World Applications.* Many decoding-phase methods face limitations in practicality, particularly in balancing time efficiency with effectiveness. Future research should prioritize practical application needs, aiming to strike an optimal balance between these factors. For example, as noted by [4], prompts remain effective in many cases, suggesting that prompt-based methods should not be overlooked. While innovative methods involving latent space manipulation and decoding-phase interventions are promising, the ultimate criterion should be their effectiveness. Researchers should select the most suitable method based on specific application scenarios to achieve the best generation outcomes.

9.2.2 Expanding the Diversity of Testing Tasks. Current testing tasks primarily focus on aspects such as toxicity, emotion, topics, and lexicon, with relatively limited evaluations of style and form. Future research should broaden the diversity of testing tasks to include aspects like linguistic style, narrative structure, and pragmatic functions. Introducing these varied testing tasks would allow for a more comprehensive evaluation of the performance and practicality of CTG models.

9.2.3 Maximizing LLM Capabilities When Comparing Baselines. When conducting experimental testing, researchers should not limit themselves to traditional CTG methods. With the advancement of LLM technology, it is essential to actively incorporate various existing prompt-based techniques to fully leverage their CTG capabilities. This approach will help in thoroughly evaluating the effectiveness of different methods, ensuring that the chosen baselines are more representative and practical, thereby identifying the optimal solution.

10 CONCLUSION

This paper reviews the latest research advances in the field of Controllable Text Generation (CTG) for Large Language Models (LLMs) and systematically defines the basic concepts, addressing both control conditions and text quality requirements. The paper introduces a new task classification approach, categorizing CTG tasks into content control (or linguistic control/hard control) and attribute control (or semantic control/soft control).

The paper provides a detailed review of various CTG methods. During the training phase, key methods include retraining or fine-tuning pre-trained models and employing reinforcement learning strategies to optimize generation quality and control precision. In the inference phase, commonly used techniques involve guiding generation through prompt engineering, manipulating the latent space for precise control, and intervening during decoding to adjust the output text.

The paper also explores various evaluation methods for CTG and highlights the wide application of CTG technology across multiple vertical domains and general tasks. The challenges faced by the CTG field, including improving quality, optimizing control precision, and enhancing inference efficiency, are discussed, along with future research directions and appeals.

In conclusion, this paper provides a comprehensive review of the core concepts, technical methods, evaluation approaches, and practical applications in the field of controllable text generation, identifying current research challenges and proposing future development directions. It aims to serve as a systematic reference and guide for research exploration in controllable text generation.

ACKNOWLEDGMENTS

This work was supported by the National Natural Science Foundation of China (Grants No. 62072463 and 71531012), the National Social Science Foundation of China (Grant No. 18ZDA309), the Research Seed Funds of the School of Interdisciplinary Studies at Renmin University of China, and the Opening Project of the State Key Laboratory of Digital Publishing Technology at Founder Group.

REFERENCES

- [1] Rohan Deepak Ajwani, Zining Zhu, Jonathan Rose, and Frank Rudzicz. 2024. Plug and Play with Prompts: A Prompt Tuning Approach for Controlling Text Generation. arXiv:2404.05143 [cs.CL] <https://arxiv.org/abs/2404.05143>
- [2] Peter Anderson, Basura Fernando, Mark Johnson, and Stephen Gould. 2016. SPICE: Semantic Propositional Image Caption Evaluation. In *Computer Vision – ECCV 2016*, Bastian Leibe, Jiri Matas, Nicu Sebe, and Max Welling (Eds.). Springer International Publishing, Cham, 382–398.
- [3] Kushal Arora, Kurt Shuster, Sainbayar Sukhbaatar, and Jason Weston. 2022. Director: Generator-Classifiers For Supervised Language Modeling. In *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, Yulan He, Heng Ji, Sujian Li, Yang Liu, and Chua-Hui Chang (Eds.). Association for Computational Linguistics, Online only, 512–526. <https://aclanthology.org/2022.aacl-main.39>

- [4] Dhananjay Ashok and Barnabas Poczos. 2024. Controllable Text Generation in the Instruction-Tuning Era. arXiv:2405.01490 [cs.CL] <https://arxiv.org/abs/2405.01490>
- [5] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Carol Chen, Catherine Olsson, Christopher Olah, Danny Hernandez, Dawn Drain, Deep Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez, Jamie Kerr, Jared Mueller, Jeffrey Ladish, Joshua Landau, Kamal Ndousse, Kamile Lukosuite, Liane Lovitt, Michael Sellitto, Nelson Elhage, Nicholas Schiefer, Noemi Mercado, Nova DasSarma, Robert Lasenby, Robin Larson, Sam Ringer, Scott Johnston, Shauna Kravec, Sheer El Showk, Stanislaw Fort, Tamera Lanham, Timothy Telleen-Lawton, Tom Conerly, Tom Henighan, Tristan Hume, Samuel R. Bowman, Zac Hatfield-Dodds, Ben Mann, Dario Amodei, Nicholas Joseph, Sam McCandlish, Tom Brown, and Jared Kaplan. 2022. Constitutional AI: Harmlessness from AI Feedback. arXiv:2212.08073 [cs.CL] <https://arxiv.org/abs/2212.08073>
- [6] Satanjeev Banerjee and Alon Lavie. 2005. METEOR: An Automatic Metric for MT Evaluation with Improved Correlation with Human Judgments. In *Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*, Jade Goldstein, Alon Lavie, Chin-Yew Lin, and Clare Voss (Eds.). Association for Computational Linguistics, Ann Arbor, Michigan, 65–72. <https://aclanthology.org/W05-0909>
- [7] Francesco Barbieri, Jose Camacho-Collados, Luis Espinosa Anke, and Leonardo Neves. 2020. TweetEval: Unified Benchmark and Comparative Evaluation for Tweet Classification. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 1644–1650. <https://doi.org/10.18653/v1/2020.findings-emnlp.148>
- [8] Emily M Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. On the dangers of stochastic parrots: Can language models be too big?. In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*. 610–623.
- [9] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language models are few-shot learners. In *Proceedings of the 34th International Conference on Neural Information Processing Systems (Vancouver, BC, Canada) (NIPS '20)*. Curran Associates Inc., Red Hook, NY, USA, Article 159, 25 pages.
- [10] Meng Cao, Mehdi Fatemi, Jackie Chi Kit Cheung, and Samira Shabanian. 2023. Successor Features for Efficient Multisubject Controlled Text Generation. arXiv:2311.04921 [cs.CL] <https://arxiv.org/abs/2311.04921>
- [11] Fredrik Carlsson, Joey Öhman, Fangyu Liu, Severine Verlinden, Joakim Nivre, and Magnus Sahlgren. 2022. Fine-Grained Controllable Text Generation Using Non-Residual Prompting. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 6837–6857. <https://doi.org/10.18653/v1/2022.acl-long.471>
- [12] Junyi Chai, Reid Pryzant, Victor Ye Dong, Konstantin Golobokov, Chenguang Zhu, and Yi Liu. 2022. FAST: Improving Controllability for Text Generation with Feedback Aware Self-Training. arXiv:2210.03167 [cs.CL] <https://arxiv.org/abs/2210.03167>
- [13] Alvin Chan, Ali Madani, Ben Krause, and Nikhil Naik. 2021. Deep Extrapolation for Attribute-Enhanced Generation. In *Advances in Neural Information Processing Systems*, A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan (Eds.). <https://openreview.net/forum?id=NCDMYD2y5kK>
- [14] Alvin Chan, Yew-Soon Ong, Bill Pung, Aston Zhang, and Jie Fu. 2021. CoCon: A Self-Supervised Approach for Controlled Text Generation. In *International Conference on Learning Representations*. https://openreview.net/forum?id=VD_ozvBy4W
- [15] Howard Chen, Huihan Li, Danqi Chen, and Karthik Narasimhan. 2022. Controllable Text Generation with Language Constraints. arXiv:2212.10466 [cs.CL] <https://arxiv.org/abs/2212.10466>
- [16] Hui Chen, Bo Wang, Ke Yang, and Yi Song. 2024. ECCRG: A Emotion- and Content-Controllable Response Generation Model. In *Collaborative Computing: Networking, Applications and Worksharing*, Honghao Gao, Xinheng Wang, and Nikolaos Voros (Eds.). Springer Nature Switzerland, Cham, 115–130.
- [17] Yihan Chen, Benfeng Xu, Quan Wang, Yi Liu, and Zhendong Mao. 2024. Benchmarking Large Language Models on Controllable Generation under Diversified Instructions. arXiv:2401.00690 [cs.CL] <https://arxiv.org/abs/2401.00690>
- [18] JinUk Cho, MinSu Jeong, JinYeong Bak, and Yun-Gyung Cheong. 2022. Genre-Controllable Story Generation via Supervised Contrastive Learning. In *Proceedings of the ACM Web Conference 2022 (Virtual Event, Lyon, France) (WWW '22)*. Association for Computing Machinery, New York, NY, USA, 2839–2849. <https://doi.org/10.1145/3485447.3512004>
- [19] cjadams, Jeffrey Sorensen, Julia Elliott, Lucas Dixon, Mark McDonald, nithum, Will Cukierski. 2018. Jigsaw Toxic Comment Classification Challenge. <https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge>

- [20] Yiming Cui, Wanxiang Che, Ting Liu, Bing Qin, Shijin Wang, and Guoping Hu. 2020. Revisiting Pre-Trained Models for Chinese Natural Language Processing. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 657–668. <https://doi.org/10.18653/v1/2020.findings-emnlp.58>
- [21] Josef Dai, Xuehai Pan, Ruiyang Sun, Jiaming Ji, Xinbo Xu, Mickel Liu, Yizhou Wang, and Yaodong Yang. 2024. Safe RLHF: Safe Reinforcement Learning from Human Feedback. In *The Twelfth International Conference on Learning Representations*. <https://openreview.net/forum?id=TyFrPOKYXw>
- [22] Sumanth Dathathri, Andrea Madotto, Janice Lan, Jane Hung, Eric Frank, Piero Molino, Jason Yosinski, and Rosanne Liu. 2020. Plug and Play Language Models: A Simple Approach to Controlled Text Generation. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=H1edEyBKDS>
- [23] Karin de Langis, Ryan Koo, and Dongyeop Kang. 2024. Reinforcement Learning with Dynamic Multi-Reward Weighting for Multi-Style Controllable Generation. arXiv:2402.14146 [cs.CL] <https://arxiv.org/abs/2402.14146>
- [24] Jasper Dekoninck, Marc Fischer, Luca Beurer-Kellner, and Martin Vechev. 2024. Controlled Text Generation via Language Model Arithmetic. In *The Twelfth International Conference on Learning Representations*. <https://openreview.net/forum?id=SLw9fp4yl6>
- [25] Haikang Deng and Colin Raffel. 2023. Reward-Augmented Decoding: Efficient Controlled Text Generation With a Unidirectional Reward Model. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 11781–11791. <https://doi.org/10.18653/v1/2023.emnlp-main.721>
- [26] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, Jill Burstein, Christy Doran, and Thamar Solorio (Eds.). Association for Computational Linguistics, Minneapolis, Minnesota, 4171–4186. <https://doi.org/10.18653/v1/N19-1423>
- [27] Hanxing Ding, Liang Pang, Zihao Wei, Huawei Shen, Xueqi Cheng, and Tat-Seng Chua. 2023. MacLaSa: Multi-Aspect Controllable Text Generation via Efficient Sampling from Compact Latent Space. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 4424–4436. <https://doi.org/10.18653/v1/2023.findings-emnlp.292>
- [28] George Doddington. 2002. Automatic evaluation of machine translation quality using n-gram co-occurrence statistics. In *Proceedings of the second international conference on Human Language Technology Research*. 138–145.
- [29] Jesse Dodge, Gabriel Ilharco, Roy Schwartz, Ali Farhadi, Hannaneh Hajishirzi, and Noah Smith. 2020. Fine-Tuning Pretrained Language Models: Weight Initializations, Data Orders, and Early Stopping. arXiv:2002.06305 [cs.CL] <https://arxiv.org/abs/2002.06305>
- [30] Chandra Kiran Reddy Evuru, Sreyan Ghosh, Sonal Kumar, Ramaneswaran S, Utkarsh Tyagi, and Dinesh Manocha. 2024. CoDa: Constrained Generation based Data Augmentation for Low-Resource NLP. arXiv:2404.00415 [cs.CL] <https://arxiv.org/abs/2404.00415>
- [31] Yuxi Feng, Xiaoyuan Yi, Xiting Wang, Laks Lakshmanan, V.S., and Xing Xie. 2023. DuNST: Dual Noisy Self Training for Semi-Supervised Controllable Text Generation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 8760–8785. <https://doi.org/10.18653/v1/2023.acl-long.488>
- [32] Zijian Feng, Hanzhang Zhou, Kezhi Mao, and Zixiao Zhu. 2024. FreeCtrl: Constructing Control Centers with Feedforward Layers for Learning-Free Controllable Text Generation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 7627–7640. <https://aclanthology.org/2024.acl-long.412>
- [33] Jarad Forristal, Fatemehsadat Mireshghallah, Greg Durrett, and Taylor Berg-Kirkpatrick. 2023. A Block Metropolis-Hastings Sampler for Controllable Energy-based Text Generation. In *Proceedings of the 27th Conference on Computational Natural Language Learning (CoNLL)*, Jing Jiang, David Reitter, and Shumin Deng (Eds.). Association for Computational Linguistics, Singapore, 403–413. <https://doi.org/10.18653/v1/2023.conll-1.26>
- [34] Samuel Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A. Smith. 2020. RealToxicityPrompts: Evaluating Neural Toxic Degeneration in Language Models. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 3356–3369. <https://doi.org/10.18653/v1/2020.findings-emnlp.301>
- [35] Yuxuan Gu, Xiaocheng Feng, Sicheng Ma, Jiaming Wu, Heng Gong, and Bing Qin. 2022. Improving Controllable Text Generation with Position-Aware Weighted Decoding. In *Findings of the Association for Computational Linguistics: ACL 2022*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 3449–3467. <https://doi.org/10.18653/v1/2022.findings-acl.272>

- [36] Yuxuan Gu, Xiaocheng Feng, Sicheng Ma, Lingyuan Zhang, Heng Gong, and Bing Qin. 2022. A Distributional Lens for Multi-Aspect Controllable Text Generation. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang (Eds.). Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 1023–1043. <https://doi.org/10.18653/v1/2022.emnlp-main.67>
- [37] Yuxuan Gu, Xiaocheng Feng, Sicheng Ma, Lingyuan Zhang, Heng Gong, Weihong Zhong, and Bing Qin. 2023. Controllable Text Generation via Probability Density Estimation in the Latent Space. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 12590–12616. <https://doi.org/10.18653/v1/2023.acl-long.704>
- [38] Bin Guo, Hao Wang, Yasan Ding, Wei Wu, Shaoyang Hao, Yueqi Sun, and Zhiwen Yu. 2021. Conditional Text Generation for Harmonious Human-Machine Interaction. *ACM Trans. Intell. Syst. Technol.* 12, 2, Article 14 (feb 2021), 50 pages. <https://doi.org/10.1145/3439816>
- [39] Xingang Guo, Fangxu Yu, Huan Zhang, Lianhui Qin, and Bin Hu. 2024. COLD-Attack: Jailbreaking LLMs with Stealthiness and Controllability. arXiv:2402.08679 [cs.LG] <https://arxiv.org/abs/2402.08679>
- [40] Skyler Hallinan, Faeze Brahman, Ximing Lu, Jaehun Jung, Sean Welleck, and Yejin Choi. 2023. STEER: Unified Style Transfer with Expert Reinforcement. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 7546–7562. <https://doi.org/10.18653/v1/2023.findings-emnlp.506>
- [41] Skyler Hallinan, Alisa Liu, Yejin Choi, and Maarten Sap. 2023. Detoxifying Text with MaRCo: Controllable Revision with Experts and Anti-Experts. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 228–242. <https://doi.org/10.18653/v1/2023.acl-short.21>
- [42] Chi Han, Jialiang Xu, Manling Li, Yi Fung, Chenkai Sun, Nan Jiang, Tarek Abdelzaher, and Heng Ji. 2024. Word Embeddings Are Steers for Language Models. arXiv:2305.12798 [cs.CL] <https://arxiv.org/abs/2305.12798>
- [43] Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. 2021. DEBERTA: DECODING-ENHANCED BERT WITH DISENTANGLED ATTENTION. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=XPZlaotutsD>
- [44] Xingwei He. 2021. Parallel Refinements for Lexically Constrained Text Generation with BART. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih (Eds.). Association for Computational Linguistics, Online and Punta Cana, Dominican Republic, 8653–8666. <https://doi.org/10.18653/v1/2021.emnlp-main.681>
- [45] Sepp Hochreiter and Jürgen Schmidhuber. 1997. Long Short-Term Memory. *Neural Comput.* 9, 8 (nov 1997), 1735–1780. <https://doi.org/10.1162/neco.1997.9.8.1735>
- [46] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. 2019. Parameter-Efficient Transfer Learning for NLP. arXiv:1902.00751 [cs.LG] <https://arxiv.org/abs/1902.00751>
- [47] Zhiting Hu and Li Erran Li. 2021. A Causal Lens for Controllable Text Generation. In *Advances in Neural Information Processing Systems*, A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan (Eds.). <https://openreview.net/forum?id=kAm9By0R5ME>
- [48] Wenyue Hua, Xianjun Yang, Mingyu Jin, Wei Cheng, Ruixiang Tang, and Yongfeng Zhang. 2024. TrustAgent: Towards Safe and Trustworthy LLM-based Agents through Agent Constitution. arXiv:2402.01586 [cs.CL] <https://arxiv.org/abs/2402.01586>
- [49] Xinyu Hua and Lu Wang. 2020. PAIR: Planning and Iterative Refinement in Pre-trained Transformers for Long Text Generation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Bonnie Webber, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 781–793. <https://doi.org/10.18653/v1/2020.emnlp-main.57>
- [50] Xuancheng Huang, Zijun Liu, Peng Li, Tao Li, Maosong Sun, and Yang Liu. 2023. An Extensible Plug-and-Play Method for Multi-Aspect Controllable Text Generation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 15233–15256. <https://doi.org/10.18653/v1/2023.acl-long.849>
- [51] Renlong Jie, Xiaojun Meng, Lifeng Shang, Xin Jiang, and Qun Liu. 2024. Prompt-Based Length Controlled Generation with Multiple Control Types. In *Findings of the Association for Computational Linguistics ACL 2024*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand and virtual meeting, 1067–1085. <https://aclanthology.org/2024.findings-acl.63>
- [52] Nie Jinran, Yang Liner, Chen Yun, Kong Cunliang, Zhu Junhui, and Yang Erhong. 2023. Lexical Complexity Controlled Sentence Generation for Language Learning. In *Proceedings of the 22nd Chinese National Conference on Computational*

- Linguistics*, Maosong Sun, Bing Qin, Xipeng Qiu, Jing Jiang, and Xianpei Han (Eds.). Chinese Information Processing Society of China, Harbin, China, 648–664. <https://aclanthology.org/2023.ccl-1.56>
- [53] Rafal Jozefowicz, Oriol Vinyals, Mike Schuster, Noam Shazeer, and Yonghui Wu. 2016. Exploring the limits of language modeling. <https://arxiv.org/pdf/1602.02410.pdf>
- [54] Juseon-Do Juseon-Do, Hidetaka Kamigaito, Manabu Okumura, and Jingun Kwon. 2024. InstructCMP: Length Control in Sentence Compression through Instruction-based Large Language Models. In *Findings of the Association for Computational Linguistics ACL 2024*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand and virtual meeting, 8980–8996. <https://aclanthology.org/2024.findings-acl.532>
- [55] Dmytro Kalpakchi and Johan Boye. 2023. SweCTRL-Mini: a data-transparent Transformer-based large language model for controllable text generation in Swedish. arXiv:2304.13994 [cs.CL] <https://arxiv.org/abs/2304.13994>
- [56] Sara Kangaslahti and David Alvarez-Melis. 2024. Continuous Language Model Interpolation for Dynamic and Controllable Text Generation. arXiv:2404.07117 [cs.CL] <https://arxiv.org/abs/2404.07117>
- [57] Pei Ke, Hao Zhou, Yankai Lin, Peng Li, Jie Zhou, Xiaoyan Zhu, and Minlie Huang. 2022. CTRLEval: An Unsupervised Reference-Free Metric for Evaluating Controlled Text Generation. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 2306–2319. <https://doi.org/10.18653/v1/2022.acl-long.164>
- [58] Nitish Shirish Keskar, Bryan McCann, Lav R. Varshney, Caiming Xiong, and Richard Socher. 2019. CTRL: A Conditional Transformer Language Model for Controllable Generation. arXiv:1909.05858 [cs.CL]
- [59] Muhammad Khalifa, Hady Elsahar, and Marc Dymetman. 2021. A Distributional Approach to Controlled Text Generation. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=jWkw45-9AbL>
- [60] Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. 2020. Generalization through Memorization: Nearest Neighbor Language Models. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=HklBjCEKvH>
- [61] Minbeom Kim, Hwanhee Lee, Kang Min Yoo, Joonsuk Park, Hwaran Lee, and Kyomin Jung. 2023. Critic-Guided Decoding for Controlled Text Generation. In *Findings of the Association for Computational Linguistics: ACL 2023*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 4598–4612. <https://doi.org/10.18653/v1/2023.findings-acl.281>
- [62] Diederik P Kingma and Max Welling. 2022. Auto-Encoding Variational Bayes. arXiv:1312.6114 [stat.ML] <https://arxiv.org/abs/1312.6114>
- [63] Tassilo Klein and Moin Nabi. 2024. Contrastive Perplexity for Controlled Generation: An Application in Detoxifying Large Language Models. arXiv:2401.08491 [cs.CL] <https://arxiv.org/abs/2401.08491>
- [64] Kai Konen, Sophie Jentzsch, Diaoulé Diallo, Peer Schütt, Oliver Bensch, Roxanne El Baff, Dominik Opitz, and Tobias Hecking. 2024. Style Vectors for Steering Generative Large Language Models. In *Findings of the Association for Computational Linguistics: EACL 2024*, Yvette Graham and Matthew Purver (Eds.). Association for Computational Linguistics, St. Julian’s, Malta, 782–802. <https://aclanthology.org/2024.findings-eacl.52>
- [65] Ben Krause, Akhilesh Deepak Gotmare, Bryan McCann, Nitish Shirish Keskar, Shafiq Joty, Richard Socher, and Nazneen Fatema Rajani. 2021. GeDi: Generative Discriminator Guided Sequence Generation. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih (Eds.). Association for Computational Linguistics, Punta Cana, Dominican Republic, 4929–4952. <https://doi.org/10.18653/v1/2021.findings-emnlp.424>
- [66] Sachin Kumar, Eric Malmi, Aliaksei Severyn, and Yulia Tsvetkov. 2021. Controlled Text Generation as Continuous Optimization with Multiple Constraints. In *Advances in Neural Information Processing Systems*, A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan (Eds.). <https://openreview.net/forum?id=kTy7bbm-4l4>
- [67] Sachin Kumar, Biswajit Paria, and Yulia Tsvetkov. 2022. Gradient-based Constrained Sampling from Language Models. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang (Eds.). Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 2251–2277. <https://doi.org/10.18653/v1/2022.emnlp-main.144>
- [68] Vaibhav Kumar, Hana Koorehdavoudi, Masud Moshtaghi, Amita Misra, Ankit Chadha, and Emilio Ferrara. 2023. Controlled Text Generation with Hidden Representation Transformations. In *Findings of the Association for Computational Linguistics: ACL 2023*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 9440–9455. <https://doi.org/10.18653/v1/2023.findings-acl.602>
- [69] Jin Myung Kwak, Minseon Kim, and Sung Ju Hwang. 2023. Language Detoxification with Attribute-Discriminative Latent Space. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 10149–10171. <https://doi.org/10.18653/v1/2023.acl-long.565>

- [70] David Landsman, Jerry Zikun Chen, and Hussain Zaidi. 2022. BeamR: Beam Reweighing with Attribute Discriminators for Controllable Text Generation. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2022*, Yulan He, Heng Ji, Sujian Li, Yang Liu, and Chua-Hui Chang (Eds.). Association for Computational Linguistics, Online only, 422–437. <https://aclanthology.org/2022.findings-acl.40>
- [71] Jean Lee, Nicholas Stevens, Soyeon Caren Han, and Minseok Song. 2024. A Survey of Large Language Models in Finance (FinLLMs). arXiv:2402.02315 [cs.CL] <https://arxiv.org/abs/2402.02315>
- [72] Chak Tou Leong, Yi Cheng, Jiahuo Wang, Jian Wang, and Wenjie Li. 2023. Self-Detoxifying Language Models via Toxicification Reversal. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 4433–4449. <https://doi.org/10.18653/v1/2023.emnlp-main.269>
- [73] Brian Lester, Rami Al-Rfou, and Noah Constant. 2021. The Power of Scale for Parameter-Efficient Prompt Tuning. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih (Eds.). Association for Computational Linguistics, Online and Punta Cana, Dominican Republic, 3045–3059. <https://doi.org/10.18653/v1/2021.emnlp-main.243>
- [74] Jiwei Li, Michel Galley, Chris Brockett, Jianfeng Gao, and Bill Dolan. 2016. A Diversity-Promoting Objective Function for Neural Conversation Models. In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, Kevin Knight, Ani Nenkova, and Owen Rambow (Eds.). Association for Computational Linguistics, San Diego, California, 110–119. <https://doi.org/10.18653/v1/N16-1014>
- [75] Wendi Li, Wei Wei, Kaihe Xu, Wenfeng Xie, Dangyang Chen, and Yu Cheng. 2024. Reinforcement Learning with Token-level Feedback for Controllable Text Generation. arXiv:2403.11558 [cs.CL] <https://arxiv.org/abs/2403.11558>
- [76] Xiang Lisa Li and Percy Liang. 2021. Prefix-Tuning: Optimizing Continuous Prompts for Generation. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, Chengqing Zong, Fei Xia, Wenjie Li, and Roberto Navigli (Eds.). Association for Computational Linguistics, Online, 4582–4597. <https://doi.org/10.18653/v1/2021.acl-long.353>
- [77] Yanran Li, Hui Su, Xiaoyu Shen, Wenjie Li, Ziqiang Cao, and Shuzi Niu. 2017. DailyDialog: A Manually Labelled Multi-turn Dialogue Dataset. In *Proceedings of the Eighth International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, Greg Kondrak and Taro Watanabe (Eds.). Asian Federation of Natural Language Processing, Taipei, Taiwan, 986–995. <https://aclanthology.org/I17-1099>
- [78] Yu Li, Zhihua Wei, Han Jiang, and Chuanyang Gong. 2024. DESTEIN: Navigating Detoxification of Language Models via Universal Steering Pairs and Head-wise Activation Fusion. arXiv:2404.10464 [cs.CL] <https://arxiv.org/abs/2404.10464>
- [79] Xun Liang, Shichao Song, Simin Niu, Zhiyu Li, Feiyu Xiong, Bo Tang, Yezhaohui Wang, Dawei He, Cheng Peng, Zhonghao Wang, and Haiying Deng. 2024. UHGEval: Benchmarking the Hallucination of Chinese Large Language Models via Unconstrained Generation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 5266–5293. <https://aclanthology.org/2024.acl-long.288>
- [80] Xun Liang, Shichao Song, Zifan Zheng, Hanyu Wang, Qingchen Yu, Xunkai Li, Rong-Hua Li, Feiyu Xiong, and Zhiyu Li. 2024. Internal Consistency and Self-Feedback in Large Language Models: A Survey. arXiv:2407.14507 [cs.CL] <https://arxiv.org/abs/2407.14507>
- [81] Xun Liang, Hanyu Wang, Shichao Song, Mengting Hu, Xunzhi Wang, Zhiyu Li, Feiyu Xiong, and Bo Tang. 2024. Controlled Text Generation for Large Language Model with Dynamic Attribute Graphs. In *Findings of the Association for Computational Linguistics ACL 2024*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand and virtual meeting, 5797–5814. <https://aclanthology.org/2024.findings-acl.345>
- [82] Bill Yuchen Lin, Abhilasha Ravichander, Ximing Lu, Nouha Dziri, Melanie Sclar, Khyathi Chandu, Chandra Bhagavatula, and Yejin Choi. 2024. The Unlocking Spell on Base LLMs: Rethinking Alignment via In-Context Learning. In *The Twelfth International Conference on Learning Representations*. <https://openreview.net/forum?id=wxJ0eXwwda>
- [83] Chin-Yew Lin. 2004. ROUGE: A Package for Automatic Evaluation of Summaries. In *Text Summarization Branches Out*. Association for Computational Linguistics, Barcelona, Spain, 74–81. <https://aclanthology.org/W04-1013>
- [84] Zhiyu Lin and Mark Riedl. 2021. Plug-and-Blend: A Framework for Controllable Story Generation with Blended Control Codes. In *Proceedings of the Third Workshop on Narrative Understanding*, Nader Akoury, Faeze Brahman, Snigdha Chaturvedi, Elizabeth Clark, Mohit Iyyer, and Lara J. Martin (Eds.). Association for Computational Linguistics, Virtual, 62–71. <https://doi.org/10.18653/v1/2021.nuse-1.7>
- [85] Alisa Liu, Maarten Sap, Ximing Lu, Swabha Swayamdipta, Chandra Bhagavatula, Noah A. Smith, and Yejin Choi. 2021. DExperts: Decoding-Time Controlled Text Generation with Experts and Anti-Experts. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, Chengqing Zong, Fei Xia, Wenjie Li, and Roberto Navigli (Eds.). Association for Computational Linguistics, Online, 6691–6706. <https://doi.org/10.18653/v1/2021.acl-long.522>

- [86] Han Liu, Bingning Wang, Ting Yao, Haijin Liang, Jianjin Xu, and Xiaolin Hu. 2022. Bridging the Gap Between Training and Inference of Bayesian Controllable Language Models. arXiv:2206.05519 [cs.CL] <https://arxiv.org/abs/2206.05519>
- [87] Sheng Liu, Haotian Ye, Lei Xing, and James Zou. 2024. In-context Vectors: Making In Context Learning More Effective and Controllable Through Latent Space Steering. arXiv:2311.06668 [cs.LG] <https://arxiv.org/abs/2311.06668>
- [88] Xin Liu, Muhammad Khalifa, and Lu Wang. 2023. BOLT: Fast Energy-based Controlled Text Generation with Tunable Biases. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 186–200. <https://doi.org/10.18653/v1/2023.acl-short.18>
- [89] Xiao Liu, Yanan Zheng, Zhengxiao Du, Ming Ding, Yujie Qian, Zhilin Yang, and Jie Tang. 2023. GPT Understands, Too. arXiv:2103.10385 [cs.CL] <https://arxiv.org/abs/2103.10385>
- [90] Yi Liu, Xiangyu Liu, Xiangrong Zhu, and Wei Hu. 2024. Multi-Aspect Controllable Text Generation with Disentangled Counterfactual Augmentation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 9231–9253. <https://aclanthology.org/2024.acl-long.500>
- [91] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692* (2019).
- [92] Yinhong Liu, Yixuan Su, Ehsan Shareghi, and Nigel Collier. 2022. Plug-and-Play Recipe Generation with Content Planning. In *Proceedings of the 2nd Workshop on Natural Language Generation, Evaluation, and Metrics (GEM)*, Antoine Bosselut, Khyathi Chandu, Kaustubh Dhole, Varun Gangal, Sebastian Gehrmann, Yacine Jernite, Jekaterina Novikova, and Laura Perez-Beltrachini (Eds.). Association for Computational Linguistics, Abu Dhabi, United Arab Emirates (Hybrid), 223–234. <https://doi.org/10.18653/v1/2022.gem-1.19>
- [93] Michela Lorandi and Anya Belz. 2023. How to Control Sentiment in Text Generation: A Survey of the State-of-the-Art in Sentiment-Control Techniques. In *Proceedings of the 13th Workshop on Computational Approaches to Subjectivity, Sentiment, & Social Media Analysis*, Jeremy Barnes, Orphée De Clercq, and Roman Klinger (Eds.). Association for Computational Linguistics, Toronto, Canada, 341–353. <https://doi.org/10.18653/v1/2023.wassa-1.30>
- [94] Ximing Lu, Sean Welleck, Peter West, Liwei Jiang, Jungo Kasai, Daniel Khashabi, Ronan Le Bras, Lianhui Qin, Youngjae Yu, Rowan Zellers, Noah A. Smith, and Yejin Choi. 2022. NeuroLogic A*esque Decoding: Constrained Text Generation with Lookahead Heuristics. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, Marine Carpuat, Marie-Catherine de Marneffe, and Ivan Vladimir Meza Ruiz (Eds.). Association for Computational Linguistics, Seattle, United States, 780–799. <https://doi.org/10.18653/v1/2022.naacl-main.57>
- [95] Ximing Lu, Peter West, Rowan Zellers, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. 2021. NeuroLogic Decoding: (Un)supervised Neural Text Generation with Predicate Logic Constraints. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, Kristina Toutanova, Anna Rumshisky, Luke Zettlemoyer, Dilek Hakkani-Tur, Iz Beltagy, Steven Bethard, Ryan Cotterell, Tanmoy Chakraborty, and Yichao Zhou (Eds.). Association for Computational Linguistics, Online, 4288–4299. <https://doi.org/10.18653/v1/2021.naacl-main.339>
- [96] Zhenyi Lu, Wei Wei, Xiaoye Qu, Xian-Ling Mao, Dangyang Chen, and Jixiong Chen. 2023. Miracle: Towards Personalized Dialogue Generation with Latent-Space Multiple Personal Attribute Control. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 5933–5957. <https://doi.org/10.18653/v1/2023.findings-emnlp.395>
- [97] Congda Ma, Tianyu Zhao, Makoto Shing, Kei Sawada, and Manabu Okumura. 2023. Focused Prefix Tuning for Controllable Text Generation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 1116–1127. <https://doi.org/10.18653/v1/2023.acl-short.96>
- [98] Andrew L. Maas, Raymond E. Daly, Peter T. Pham, Dan Huang, Andrew Y. Ng, and Christopher Potts. 2011. Learning Word Vectors for Sentiment Analysis. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, Dekang Lin, Yuji Matsumoto, and Rada Mihalcea (Eds.). Association for Computational Linguistics, Portland, Oregon, USA, 142–150. <https://aclanthology.org/P11-1015>
- [99] Julian McAuley and Jure Leskovec. 2013. Hidden factors and hidden topics: understanding rating dimensions with review text. In *Proceedings of the 7th ACM conference on Recommender systems*. 165–172.
- [100] Tao Meng, Sidi Lu, Nanyun Peng, and Kai-Wei Chang. 2022. Controllable Text Generation with Neurally-Decomposed Oracle. In *Advances in Neural Information Processing Systems*, S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, and A. Oh (Eds.), Vol. 35. Curran Associates, Inc., 28125–28139. https://proceedings.neurips.cc/paper_files/paper/2022/file/b40d5797756800c97f3d525c2e4c8357-Paper-Conference.pdf

- [101] Fatemehsadat Mireshghallah, Kartik Goyal, and Taylor Berg-Kirkpatrick. 2022. Mix and Match: Learning-free Controllable Text Generation using Energy Language Models. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 401–415. <https://doi.org/10.18653/v1/2022.acl-long.31>
- [102] Sidharth Mudgal, Jong Lee, Harish Ganapathy, YaGuang Li, Tao Wang, Yanping Huang, Zhifeng Chen, Heng-Tze Cheng, Michael Collins, Jilin Chen, Alex Beutel, and Ahmad Beirami. 2023. Controlled Decoding from Language Models. In *Socially Responsible Language Modelling Research*. <https://openreview.net/forum?id=jo57H1CpD8>
- [103] Gilles Nawezi, Lucie Flek, and Charles Welch. 2023. Style Locality for Controllable Generation with kNN Language Models. In *Proceedings of the 1st Workshop on Taming Large Language Models: Controllability in the era of Interactive Assistants!*, Devamanyu Hazarika, Xiangru Robert Tang, and Di Jin (Eds.). Association for Computational Linguistics, Prague, Czech Republic, 68–75. <https://aclanthology.org/2023.tllm-1.7>
- [104] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F Christiano, Jan Leike, and Ryan Lowe. 2022. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, and A. Oh (Eds.), Vol. 35. Curran Associates, Inc., 27730–27744. https://proceedings.neurips.cc/paper_files/paper/2022/file/b1efde53be364a73914f58805a001731-Paper-Conference.pdf
- [105] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. Bleu: a Method for Automatic Evaluation of Machine Translation. In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, Pierre Isabelle, Eugene Charniak, and Dekang Lin (Eds.). Association for Computational Linguistics, Philadelphia, Pennsylvania, USA, 311–318. <https://doi.org/10.3115/1073083.1073135>
- [106] Damian Pascual, Beni Egressy, Clara Meister, Ryan Cotterell, and Roger Wattenhofer. 2021. A Plug-and-Play Method for Controlled Text Generation. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih (Eds.). Association for Computational Linguistics, Punta Cana, Dominican Republic, 3973–3997. <https://doi.org/10.18653/v1/2021.findings-emnlp.334>
- [107] Jonathan Pei, Kevin Yang, and Dan Klein. 2023. PREADD: Prefix-Adaptive Decoding for Controlled Text Generation. In *Findings of the Association for Computational Linguistics: ACL 2023*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 10018–10037. <https://doi.org/10.18653/v1/2023.findings-acl.636>
- [108] Luiza Pozzobon, Beyza Ermis, Patrick Lewis, and Sara Hooker. 2023. Goodtriever: Adaptive Toxicity Mitigation with Retrieval-augmented Models. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 5108–5125. <https://doi.org/10.18653/v1/2023.findings-emnlp.339>
- [109] Shrimai Prabhumoye, Alan W Black, and Ruslan Salakhutdinov. 2020. Exploring Controllable Text Generation Techniques. In *Proceedings of the 28th International Conference on Computational Linguistics*, Donia Scott, Nuria Bel, and Chengqing Zong (Eds.). International Committee on Computational Linguistics, Barcelona, Spain (Online), 1–14. <https://doi.org/10.18653/v1/2020.coling-main.1>
- [110] Jing Qian, Li Dong, Yelong Shen, Furu Wei, and Weizhu Chen. 2022. Controllable Natural Language Generation with Contrastive Prefixes. In *Findings of the Association for Computational Linguistics: ACL 2022*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 2912–2924. <https://doi.org/10.18653/v1/2022.findings-acl.229>
- [111] Lianhui Qin, Sean Welleck, Daniel Khashabi, and Yejin Choi. 2022. COLD Decoding: Energy-based Constrained Text Generation with Langevin Dynamics. In *Advances in Neural Information Processing Systems*, S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, and A. Oh (Eds.), Vol. 35. Curran Associates, Inc., 9538–9551. https://proceedings.neurips.cc/paper_files/paper/2022/file/3e25d1aff47964c8409fd5c8dc0438d7-Paper-Conference.pdf
- [112] Alec Radford, Luke Metz, and Soumith Chintala. 2016. Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks. In *4th International Conference on Learning Representations, ICLR 2016, San Juan, Puerto Rico, May 2–4, 2016, Conference Track Proceedings*.
- [113] Abir Rahali and Moulay A. Akhloufi. 2023. DeepPress: guided press release topic-aware text generation using ensemble transformers. *Neural Computing and Applications* 35, 17 (2023), 12847–12874. <https://doi.org/10.1007/s00521-023-08393-4>
- [114] Angela Ramirez, Kartik Agarwal, Juraj Juraska, Utkarsh Garg, and Marilyn Walker. 2023. Controllable Generation of Dialogue Acts for Dialogue Systems via Few-Shot Response Generation and Ranking. In *Proceedings of the 24th Annual Meeting of the Special Interest Group on Discourse and Dialogue*, Svetlana Stoyanchev, Shafiq Joty, David Schlangen, Ondrej Dusek, Casey Kennington, and Malihe Alikhani (Eds.). Association for Computational Linguistics, Prague, Czechia, 355–369. <https://doi.org/10.18653/v1/2023.sigdial-1.32>

- [115] Marc'Aurelio Ranzato, Sumit Chopra, Michael Auli, and Wojciech Zaremba. 2016. Sequence Level Training with Recurrent Neural Networks. arXiv:1511.06732 [cs.LG] <https://arxiv.org/abs/1511.06732>
- [116] David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. 1986. Learning representations by back-propagating errors. *Nature* 323, 6088 (1986), 533–536. <https://doi.org/10.1038/323533a0>
- [117] Punyajoy Saha, Kanishk Singh, Adarsh Kumar, Binny Mathew, and Animesh Mukherjee. 2022. CounterGeDi: A Controllable Approach to Generate Polite, Detoxified and Emotional Counterspeech. In *Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence, IJCAI-22*, Lud De Raedt (Ed.). International Joint Conferences on Artificial Intelligence Organization, 5157–5163. <https://doi.org/10.24963/ijcai.2022/716> AI for Good.
- [118] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. 2019. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108* (2019).
- [119] Yuichi Sasazawa, Terufumi Morishita, Hiroaki Ozaki, Osamu Imaichi, and Yasuhiro Sogawa. 2023. Controlling keywords and their positions in text generation. In *Proceedings of the 16th International Natural Language Generation Conference*, C. Maria Keet, Hung-Yi Lee, and Sina Zarrieß (Eds.). Association for Computational Linguistics, Prague, Czechia, 407–413. <https://doi.org/10.18653/v1/2023.inlg-main.29>
- [120] Timo Schick, Sahana Udupa, and Hinrich Schütze. 2021. Self-Diagnosis and Self-Debiasing: A Proposal for Reducing Corpus-Based Bias in NLP. *Transactions of the Association for Computational Linguistics* 9 (2021), 1408–1424. https://doi.org/10.1162/tacl_a_00434
- [121] Thibault Sellam, Dipanjan Das, and Ankur Parikh. 2020. BLEURT: Learning Robust Metrics for Text Generation. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel Tetreault (Eds.). Association for Computational Linguistics, Online, 7881–7892. <https://doi.org/10.18653/v1/2020.acl-main.704>
- [122] Damith Chamalke Senadeera and Julia Ive. 2022. Controlled Text Generation using T5 based Encoder-Decoder Soft Prompt Tuning and Analysis of the Utility of Generated Text in AI. arXiv:2212.02924 [cs.CL] <https://arxiv.org/abs/2212.02924>
- [123] Zhihong Shao, Minlie Huang, Jiangtao Wen, Wenfei Xu, and Xiaoyan Zhu. 2019. Long and Diverse Text Generation with Planning-based Hierarchical Variational Model. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan (Eds.). Association for Computational Linguistics, Hong Kong, China, 3257–3268. <https://doi.org/10.18653/v1/D19-1321>
- [124] Chenhui Shen, Liying Cheng, Ran Zhou, Lidong Bing, Yang You, and Luo Si. 2022. MRd: A Meta-Review Dataset for Structure-Controllable Text Generation. In *Findings of the Association for Computational Linguistics: ACL 2022*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 2521–2535. <https://doi.org/10.18653/v1/2022.findings-acl.198>
- [125] Chufan Shi, Deng Cai, and Yujiu Yang. 2024. LiFi: Lightweight Controlled Text Generation with Fine-Grained Control Codes. arXiv:2402.06930 [cs.CL] <https://arxiv.org/abs/2402.06930>
- [126] Taylor Shin, Yasaman Razeghi, Robert L. Logan IV, Eric Wallace, and Sameer Singh. 2020. AutoPrompt: Eliciting Knowledge from Language Models with Automatically Generated Prompts. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Bonnie Webber, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 4222–4235. <https://doi.org/10.18653/v1/2020.emnlp-main.346>
- [127] Askhat Sitdikov, Nikita Balagansky, Daniil Gavrillov, and Alexander Markov. 2022. Classifiers are Better Experts for Controllable Text Generation. arXiv:2205.07276 [cs.CL] <https://arxiv.org/abs/2205.07276>
- [128] Matthew Snover, Bonnie Dorr, Rich Schwartz, Linnea Micciulla, and John Makhoul. 2006. A Study of Translation Edit Rate with Targeted Human Annotation. In *Proceedings of the 7th Conference of the Association for Machine Translation in the Americas: Technical Papers*. Association for Machine Translation in the Americas, Cambridge, Massachusetts, USA, 223–231. <https://aclanthology.org/2006.amta-papers.25>
- [129] Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Ng, and Christopher Potts. 2013. Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, David Yarowsky, Timothy Baldwin, Anna Korhonen, Karen Livescu, and Steven Bethard (Eds.). Association for Computational Linguistics, Seattle, Washington, USA, 1631–1642. <https://aclanthology.org/D13-1170>
- [130] Alexander Spangher, Yao Ming, Xinyu Hua, and Nanyun Peng. 2022. Sequentially Controlled Text Generation. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang (Eds.). Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 6848–6866. <https://doi.org/10.18653/v1/2022.findings-emnlp.509>
- [131] Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. 2020. Learning to summarize with human feedback. In *Advances in Neural Information Processing*

- Systems, H. Larochelle, M. Ranzato, R. Hadsell, M.F. Balcan, and H. Lin (Eds.), Vol. 33. Curran Associates, Inc., 3008–3021. https://proceedings.neurips.cc/paper_files/paper/2020/file/1f89885d556929e98d3ef9b86448f951-Paper.pdf
- [132] Nishant Subramani, Nivedita Suresh, and Matthew Peters. 2022. Extracting Latent Steering Vectors from Pretrained Language Models. In *Findings of the Association for Computational Linguistics: ACL 2022*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 566–581. <https://doi.org/10.18653/v1/2022.findings-acl.48>
- [133] Richard S. Sutton, David McAllester, Satinder Singh, and Yishay Mansour. 1999. Policy gradient methods for reinforcement learning with function approximation. In *Proceedings of the 12th International Conference on Neural Information Processing Systems* (Denver, CO) (NIPS’99). MIT Press, Cambridge, MA, USA, 1057–1063.
- [134] Zhen Tao, Dinghao Xi, Zhiyu Li, Liumin Tang, and Wei Xu. 2024. CAT-LLM: Prompting Large Language Models with Text Style Definition for Chinese Article-style Transfer. arXiv:2401.05707 [cs.CL] <https://arxiv.org/abs/2401.05707>
- [135] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. 2023. LLaMA: Open and Efficient Foundation Language Models. arXiv:2302.13971 [cs.CL] <https://arxiv.org/abs/2302.13971>
- [136] Severino Trotta, Lucie Flek, and Charles Welch. 2022. Nearest Neighbor Language Models for Stylistic Controllable Generation. In *Proceedings of the 2nd Workshop on Natural Language Generation, Evaluation, and Metrics (GEM)*, Antoine Bosselut, Khyathi Chandu, Kaustubh Dhole, Varun Gangal, Sebastian Gehrmann, Yacine Jernite, Jekaterina Novikova, and Laura Perez-Beltrachini (Eds.). Association for Computational Linguistics, Abu Dhabi, United Arab Emirates (Hybrid), 295–305. <https://doi.org/10.18653/v1/2022.gem-1.25>
- [137] Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udeh, Juan J. Vazquez, Ulisse Mini, and Monte MacDiarmid. 2024. Activation Addition: Steering Language Models Without Optimization. arXiv:2308.10248 [cs.CL] <https://arxiv.org/abs/2308.10248>
- [138] Bhargav Upadhyay, Akhilesh Sudhakar, and Arjun Maheswaran. 2022. Efficient Reinforcement Learning for Unsupervised Controlled Text Generation. arXiv:2204.07696 [cs.CL] <https://arxiv.org/abs/2204.07696>
- [139] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is All you Need. In *Advances in Neural Information Processing Systems*, I. Guyon, U. Von Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett (Eds.), Vol. 30. Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf
- [140] Ramakrishna Vedantam, C. Lawrence Zitnick, and Devi Parikh. 2015. CIDER: Consensus-Based Image Description Evaluation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
- [141] Sergey Vychezhzhanin, Anastasia Kotelnikova, Alexander Sergeev, and Evgeny Kotelnikov. 2024. Controllable Story Generation Based on Perplexity Minimization. In *Analysis of Images, Social Networks and Texts*, Dmitry I. Ignatov, Michael Khachay, Andrey Kutuzov, Habet Madoyan, Ilya Makarov, Irina Nikishina, Alexander Panchenko, Maxim Panov, Panos M. Pardalos, Andrey V. Savchenko, Evgenii Tsybmalov, Elena Tutubalina, and Sergey Zagoruyko (Eds.). Springer Nature Switzerland, Cham, 154–169.
- [142] S. V. Vychezhzhanin and E. V. Kotelnikov. 2022. Collocation2Text: Controllable Text Generation from Guide Phrases in Russian. In *Computational Linguistics and Intellectual Technologies*. RSUH. <https://doi.org/10.28995/2075-7182-2022-21-564-576>
- [143] Zhongwei Wan, Xin Wang, Che Liu, Samiul Alam, Yu Zheng, et al. 2023. Efficient large language models: A survey. *arXiv preprint arXiv:2312.03863* 1 (2023).
- [144] Hao Wang and Lei Sha. 2023. Harnessing the Plug-and-Play Controller by Prompting. In *Proceedings of the Third Workshop on Natural Language Generation, Evaluation, and Metrics (GEM)*, Sebastian Gehrmann, Alex Wang, João Sedoc, Elizabeth Clark, Kaustubh Dhole, Khyathi Raghavi Chandu, Enrico Santus, and Hooman Sedghamiz (Eds.). Association for Computational Linguistics, Singapore, 165–174. <https://aclanthology.org/2023.gem-1.14>
- [145] Junli Wang, Chenyang Zhang, Dongyu Zhang, Haibo Tong, Chungang Yan, and Changjun Jiang. 2024. A Recent Survey on Controllable Text Generation: a Causal Perspective. *Fundamental Research* (2024). <https://api.semanticscholar.org/CorpusID:266926474>
- [146] Pengyu Wang, Dong Zhang, Linyang Li, Chenkun Tan, Xinghao Wang, Ke Ren, Botian Jiang, and Xipeng Qiu. 2024. InferAligner: Inference-Time Alignment for Harmlessness through Cross-Model Guidance. arXiv:2401.11206 [cs.CL] <https://arxiv.org/abs/2401.11206>
- [147] Xinpeng Wang, Han Jiang, Zhihua Wei, and Shanlin Zhou. 2022. CHAE: Fine-Grained Controllable Story Generation with Characters, Actions and Emotions. In *Proceedings of the 29th International Conference on Computational Linguistics*, Nicoletta Calzolari, Chu-Ren Huang, Hansaem Kim, James Pustejovsky, Leo Wanner, Key-Sun Choi, Pum-Mo Ryu, Hsin-Hsi Chen, Lucia Donatelli, Heng Ji, Sadao Kurohashi, Patrizia Paggio, Nianwen Xue, Seokhwan Kim, Younggyun Hahm, Zhong He, Tony Kyungil Lee, Enrico Santus, Francis Bond, and Seung-Hoon Na (Eds.). International Committee on Computational Linguistics, Gyeongju, Republic of Korea, 6426–6435. <https://aclanthology.org/2022.coling-1.559>

- [148] Jason Wei, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V Le. 2022. Finetuned Language Models are Zero-Shot Learners. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=gEZrGCzodqR>
- [149] Zhihua Wen, Zhiliang Tian, Zhen Huang, Yuxin Yang, Zexin Jian, Changjian Wang, and Dongsheng Li. 2023. GRACE: Gradient-guided Controllable Retrieval for Augmenting Attribute-based Text Generation. In *Findings of the Association for Computational Linguistics: ACL 2023*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 8377–8398. <https://doi.org/10.18653/v1/2023.findings-acl.530>
- [150] Congying Xia, Chen Xing, Jiangshu Du, Xinyi Yang, Yihao Feng, Ran Xu, Wenpeng Yin, and Caiming Xiong. 2024. FOFO: A Benchmark to Evaluate LLMs' Format-Following Capability. *arXiv preprint arXiv:2402.18667* (2024).
- [151] Chen Xu, Tian Lan, Changlong Yu, Wei Wang, Jun Gao, Yu Ji, Qunxi Dong, Kun Qian, Piji Li, Wei Bi, and Bin Hu. 2024. DECIDER: A Dual-System Rule-Controllable Decoding Framework for Language Generation. *arXiv:2403.01954* [cs.CL] <https://arxiv.org/abs/2403.01954>
- [152] Peng Xu, Mostofa Patwary, Mohammad Shoneybi, Raul Puri, Pascale Fung, Anima Anandkumar, and Bryan Catanzaro. 2020. MEGATRON-CNTRL: Controllable Story Generation with External Knowledge Using Large-Scale Language Models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Bonnie Webber, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 2831–2845. <https://doi.org/10.18653/v1/2020.emnlp-main.226>
- [153] Kevin Yang and Dan Klein. 2021. FUDGE: Controlled Text Generation With Future Discriminators. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, Kristina Toutanova, Anna Rumshisky, Luke Zettlemoyer, Dilek Hakkani-Tur, Iz Beltagy, Steven Bethard, Ryan Cotterell, Tanmoy Chakraborty, and Yichao Zhou (Eds.). Association for Computational Linguistics, Online, 3511–3535. <https://doi.org/10.18653/v1/2021.naacl-main.276>
- [154] Kexin Yang, Dayiheng Liu, Wenqiang Lei, Baosong Yang, Mingfeng Xue, Boxing Chen, and Jun Xie. 2023. Tailor: A Soft-Prompt-Based Approach to Attribute-Based Controlled Text Generation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 410–427. <https://doi.org/10.18653/v1/2023.acl-long.25>
- [155] Zhe Yang, Yi Huang, Yaqin Chen, Xiaoting Wu, Junlan Feng, and Chao Deng. 2024. CTGGAN: Controllable Text Generation with Generative Adversarial Network. *Applied Sciences* 14, 7 (2024). <https://doi.org/10.3390/app14073106>
- [156] Zhian Yang, Hao Jiang, Aobo Deng, and Yang Li. 2024. Topic-Oriented Controlled Text Generation for Social Networks. *Journal of Signal Processing Systems* 96, 2 (feb 2024), 131–151. <https://doi.org/10.1007/s11265-023-01907-2>
- [157] Dian Yu, Zhou Yu, and Kenji Sagae. 2021. Attribute Alignment: Controlling Text Generation from Pre-trained Language Models. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih (Eds.). Association for Computational Linguistics, Punta Cana, Dominican Republic, 2251–2268. <https://doi.org/10.18653/v1/2021.findings-emnlp.194>
- [158] Lantao Yu, Weinan Zhang, Jun Wang, and Yong Yu. 2017. SeqGAN: Sequence Generative Adversarial Nets with Policy Gradient. In *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence*. 2852–2858. <https://ojs.aaai.org/index.php/AAAI/article/view/10770>
- [159] Sangwon Yu, Changmin Lee, Hoin Lee, and Sungroh Yoon. 2024. Controlled Text Generation for Black-box Language Models via Score-based Progressive Editor. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 14215–14237. <https://aclanthology.org/2024.acl-long.767>
- [160] Yoel Zeldes, Dan Padnos, Or Sharir, and Barak Peleg. 2020. Technical Report: Auxiliary Tuning and its Application to Conditional Text Generation. *arXiv:2006.16823* [cs.CL] <https://arxiv.org/abs/2006.16823>
- [161] Qingcheng Zeng, Mingyu Jin, Qinkai Yu, Zhenting Wang, Wenyue Hua, Zihao Zhou, Guangyan Sun, Yanda Meng, Shiqing Ma, Qifan Wang, Felix Juefei-Xu, Kaize Ding, Fan Yang, Ruixiang Tang, and Yongfeng Zhang. 2024. Uncertainty is Fragile: Manipulating Uncertainty in Large Language Models. *arXiv:2407.11282* [cs.CL] <https://arxiv.org/abs/2407.11282>
- [162] Weihao Zeng, Lulu Zhao, Keqing He, Ruotong Geng, Jingang Wang, Wei Wu, and Weiran Xu. 2023. Seen to Unseen: Exploring Compositional Generalization of Multi-Attribute Controllable Dialogue Generation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 14179–14196. <https://doi.org/10.18653/v1/2023.acl-long.793>
- [163] Yongcheng Zeng, Guoqing Liu, Weiylu Ma, Ning Yang, Haifeng Zhang, and Jun Wang. 2024. Token-level Direct Preference Optimization. *arXiv:2404.11999* [cs.CL] <https://arxiv.org/abs/2404.11999>
- [164] Yan Zeng and Jian-Yun Nie. 2021. A Simple and Efficient Multi-Task Learning Approach for Conditioned Dialogue Generation. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational*

- Linguistics: Human Language Technologies*, Kristina Toutanova, Anna Rumshisky, Luke Zettlemoyer, Dilek Hakkani-Tur, Iz Beltagy, Steven Bethard, Ryan Cotterell, Tanmoy Chakraborty, and Yichao Zhou (Eds.). Association for Computational Linguistics, Online, 4927–4939. <https://doi.org/10.18653/v1/2021.naacl-main.392>
- [165] Hanqing Zhang and Dawei Song. 2022. DisCup: Discriminator Cooperative Unlikelihood Prompt-tuning for Controllable Text Generation. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang (Eds.). Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 3392–3406. <https://doi.org/10.18653/v1/2022.emnlp-main.223>
- [166] Hanqing Zhang, Haolin Song, Shaoyu Li, Ming Zhou, and Dawei Song. 2023. A Survey of Controllable Text Generation Using Transformer-based Pre-trained Language Models. *ACM Comput. Surv.* 56, 3, Article 64 (oct 2023), 37 pages. <https://doi.org/10.1145/3617680>
- [167] Hanqing Zhang, Si Sun, Haiming Wu, and Dawei Song. 2024. Controllable Text Generation with Residual Memory Transformer. In *Findings of the Association for Computational Linguistics ACL 2024*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand and virtual meeting, 1048–1066. <https://aclanthology.org/2024.findings-acl.62>
- [168] Jingyu Zhang, James Glass, and Tianxing He. 2023. PCFG-Based Natural Language Interface Improves Generalization for Controlled Text Generation. In *Proceedings of the 12th Joint Conference on Lexical and Computational Semantics (*SEM 2023)*, Alexis Palmer and Jose Camacho-collados (Eds.). Association for Computational Linguistics, Toronto, Canada, 295–313. <https://doi.org/10.18653/v1/2023.starsem-1.27>
- [169] Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. BERTScore: Evaluating Text Generation with BERT. In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26–30, 2020*. OpenReview.net. <https://openreview.net/forum?id=SkeHuCVFDr>
- [170] Xu Zhang and Xiaojun Wan. 2023. MIL-Decoding: Detoxifying Language Models at Token-Level via Multiple Instance Learning. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 190–202. <https://doi.org/10.18653/v1/2023.acl-long.11>
- [171] Xiang Zhang, Junbo Zhao, and Yann LeCun. 2015. Character-level convolutional networks for text classification. *Advances in neural information processing systems* 28 (2015).
- [172] Yizhe Zhang, Guoyin Wang, Chunyuan Li, Zhe Gan, Chris Brockett, and Bill Dolan. 2020. POINTER: Constrained Progressive Text Generation via Insertion-based Generative Pre-training. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Bonnie Webber, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 8649–8670. <https://doi.org/10.18653/v1/2020.emnlp-main.698>
- [173] Zhiling Zhang, Mengyue Wu, and Kenny Zhu. 2023. Semantic Space Grounded Weighted Decoding for Multi-Attribute Controllable Dialogue Generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 13230–13243. <https://doi.org/10.18653/v1/2023.emnlp-main.817>
- [174] Wei Zhao, Maxime Peyrard, Fei Liu, Yang Gao, Christian M. Meyer, and Steffen Eger. 2019. MoverScore: Text Generation Evaluating with Contextualized Embeddings and Earth Mover Distance. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan (Eds.). Association for Computational Linguistics, Hong Kong, China, 563–578. <https://doi.org/10.18653/v1/D19-1053>
- [175] Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Xiaolei Wang, Yupeng Hou, Yingqian Min, Beichen Zhang, Junjie Zhang, Zican Dong, Yifan Du, Chen Yang, Yushuo Chen, Zhipeng Chen, Jinhao Jiang, Ruiyang Ren, Yifan Li, Xinyu Tang, Zikang Liu, Peiyu Liu, Jian-Yun Nie, and Ji-Rong Wen. 2023. A Survey of Large Language Models. arXiv:2303.18223 [cs.CL] <https://arxiv.org/abs/2303.18223>
- [176] Chujie Zheng, Pei Ke, Zheng Zhang, and Minlie Huang. 2023. Click: Controllable Text Generation with Sequence Likelihood Contrastive Learning. In *Findings of the Association for Computational Linguistics: ACL 2023*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 1022–1040. <https://doi.org/10.18653/v1/2023.findings-acl.65>
- [177] Carolina Zheng, Claudia Shi, Keyon Vafa, Amir Feder, and David Blei. 2023. An Invariant Learning Characterization of Controlled Text Generation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 3186–3206. <https://doi.org/10.18653/v1/2023.acl-long.179>
- [178] Xin Zheng, Hongyu Lin, Xianpei Han, and Le Sun. 2023. Toward Unified Controllable Text Generation via Regular Expression Instruction. In *Proceedings of the 13th International Joint Conference on Natural Language Processing and the 3rd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, Jong C. Park, Yuki Arase, Baotian Hu, Wei Lu, Derry Wijaya, Ayu Purwarianti, and Adila Alfa Krisnadhi (Eds.).

- Association for Computational Linguistics, Nusa Dua, Bali, 1–14. <https://doi.org/10.18653/v1/2023.ijcnlp-main.1>
- [179] Yinhe Zheng, Rongsheng Zhang, Minlie Huang, and Xiaoxi Mao. 2020. A Pre-Training Based Personalized Dialogue Generation Model with Persona-Sparse Data. *Proceedings of the AAAI Conference on Artificial Intelligence* 34, 05 (Apr. 2020), 9693–9700. <https://doi.org/10.1609/aaai.v34i05.6518>
- [180] Qihuang Zhong, Liang Ding, Juhua Liu, Bo Du, and Dacheng Tao. 2024. ROSE Doesn't Do That: Boosting the Safety of Instruction-Tuned Large Language Models with Reverse Prompt Contrastive Decoding. arXiv:2402.11889 [cs.CL] <https://arxiv.org/abs/2402.11889>
- [181] Tianqi Zhong, Quan Wang, Jingxuan Han, Yongdong Zhang, and Zhendong Mao. 2023. Air-Decoding: Attribute Distribution Reconstruction for Decoding-Time Controllable Text Generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 8233–8248. <https://doi.org/10.18653/v1/2023.emnlp-main.512>
- [182] Chunting Zhou, Pengfei Liu, Puxin Xu, Srini Iyer, Jiao Sun, Yuning Mao, Xuezhe Ma, Avia Efrat, Ping Yu, LILI YU, Susan Zhang, Gargi Ghosh, Mike Lewis, Luke Zettlemoyer, and Omer Levy. 2023. LIMA: Less Is More for Alignment. In *Thirty-seventh Conference on Neural Information Processing Systems*. <https://openreview.net/forum?id=KBMOKmX2he>
- [183] Wangchunshu Zhou, Yuchen Eleanor Jiang, Ethan Wilcox, Ryan Cotterell, and Mrinmaya Sachan. 2023. Controlled text generation with natural language instructions. In *Proceedings of the 40th International Conference on Machine Learning* (Honolulu, Hawaii, USA) (ICML '23). JMLR.org, Article 1795, 12 pages.
- [184] Linan Zhu, Yifei Xu, Zhechao Zhu, Yinwei Bao, and Xiangjie Kong. 2023. Fine-Grained Sentiment-Controlled Text Generation Approach Based on Pre-Trained Language Model. *Applied Sciences* 13, 1 (2023). <https://doi.org/10.3390/app13010264>
- [185] Yaoming Zhu, Sidi Lu, Lei Zheng, Jiaxian Guo, Weinan Zhang, Jun Wang, and Yong Yu. 2018. Texygen: A benchmarking platform for text generation models. In *The 41st international ACM SIGIR conference on research & development in information retrieval*. 1097–1100.
- [186] Xu Zou, Da Yin, Qingyang Zhong, Hongxia Yang, Zhilin Yang, and Jie Tang. 2021. Controllable Generation from Pre-trained Language Models via Inverse Prompting. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining* (Virtual Event, Singapore) (KDD '21). Association for Computing Machinery, New York, NY, USA, 2450–2460. <https://doi.org/10.1145/3447548.3467418>

Received XX XX XXXX; revised XX XX XXXX; accepted XX XX XXXX